

# FIT3143

# Parallel Computing

Week 3: Inter-Process Communications and RPC

Vishnu Monn and ABM Russel



[www.shutterstock.com](http://www.shutterstock.com) · 1062915266

# Unit Topics

Week	Lecture Topic	Tutorial	Laboratory	Remarks
1	Introduction to Parallel Computing and Distributed Systems	None	Lab Week 01 - Introduction to Linux & Setting up VM (Not assessed)	Group formation for lab activities and assignments (Two students per group)
2	Parallel computing on shared memory (POSIX and OpenMP)	Tutorial Week 02 - Introduction to Parallel Computing	Lab Week 02 - C Primer (Not assessed)	
3	Inter Process Communications; Remote Procedure Calls	Tutorial Week 03 - Shared memory parallelism	Lab Week 03 - Threads (POSIX)	Assignment 1 specifications released
4	Parallel computing on distributed memory (MPI)	Tutorial Week 04 - Inter process communication	Lab Week 04 - Threads (OpenMP)	
5	Synchronization, MUTEX, Deadlocks	Tutorial Week 05 - Distributed memory parallelism	Lab Week 05 - Message passing interface	Assignment 2 specifications released
6	Election Algorithms, Distributed Transactions, Concurrency Control	Tutorial Week 06 - Synchronization	Lab Week 06 - Communication patterns	
7	Faults, Distributed Consensus, Security, Parallel Computing	Tutorial Week 07 - Transactions and concurrency	Lab Week 07 - Parallel data structure (I)	Assignment 1 due
8	Instruction Level Parallelism	Tutorial Week 08 - Consensus	1st assignment interview	
9	Vector Architecture	Tutorial Week 09 - Instruction level parallelism	Lab Week 09 - Parallel data structure (II)	
10	Data Parallel Architectures, SIMD Architectures	Tutorial Week 10 - Vector architecture	Lab Week 10 - Master slave	
11	Introduction to MIMD, Distributed Memory MIMD Architectures	Tutorial Week 11 - Data parallelism	Lab Week 11 - Performance tuning	Assignment 2 due
12	Recent development in Parallel Computing - Software, Hardware, Applications etc. (GPGPU etc.), Exam Revision	Tutorial Week 12 - Revision	2nd assignment code demo and interview	

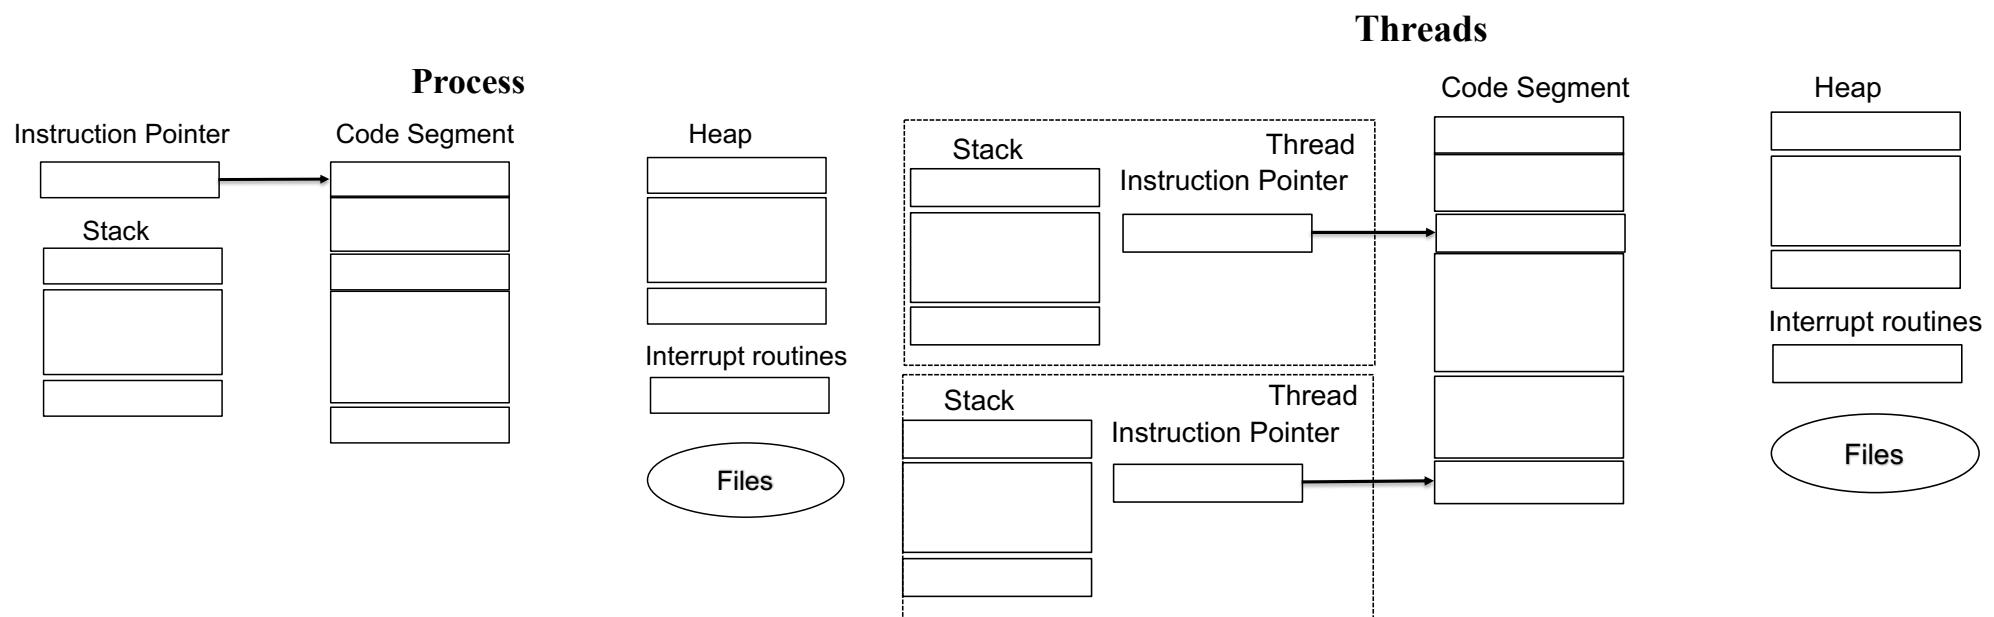
Assessment	Weight
Assignment 1	15%
Assignment 2	25%
Lab-work and Quizzes	10%
Final exam	50%

# Today

- Recap: POSIX Threads and OpenMP
- Inter-Process Communications
- RPC

# Recap

# Process vs Threads



# FLUX: Process vs Threads

**Process vs Threads**

**JYGAFY**

# POSIX Threads

```
#include <stdio.h>
#include <pthread.h>
#define NUM_THREADS 4

void *hello (void *arg) {
    printf("Hello Thread\n");
}

main() {
    pthread_t tid[NUM_THREADS];
    for (int i = 0; i < NUM_THREADS; i++)
        pthread_create(&tid[i], NULL, hello, NULL);

    for (int i = 0; i < NUM_THREADS; i++)
        pthread_join(tid[i], NULL);
}
```

# Thread rank

```
void *threadFunc(void *pArg)
{
    int myNum = *((int*)pArg);
    printf( "Thread number %d\n", myNum);
}

...
// from main():
for (int i = 0; i < numThreads; i++) {
    tNum[i] = i;
    pthread_create(&tid[i], NULL, threadFunc, &tNum[i]);
}
```

# OpenMP

## Serial Code

```
#include <stdio.h>

int main() {
    int i;
    printf("Hello World\n");

    for(i=0;i<6;i++)
        printf("Iter:%d\n",i);

    printf("GoodBye World\n");
}
```

## Parallel Code

```
#include <stdio.h>
#include <omp.h>

int main() {
#pragma omp parallel
{
    int i;
    printf("Hello World\n");

#pragma omp for
    for(i=0;i<6;i++)
        printf("Iter:%d\n",i);

}
    printf("GoodBye World\n");
}
```

# The Private & Shared Clause

```
void* work(float* c, int N) {  
    float x, y; int i;  
#pragma omp parallel for private(x,y)  
    for(i=0; i<N; i++) {  
        x = a[i]; y = b[i];  
        c[i] = x + y;  
    }  
}
```

```
float dot_prod(float* a, float* b, int N)  
{  
    float sum = 0.0;  
#pragma omp parallel for shared(sum)  
    for(int i=0; i<N; i++) {  
        sum += a[i] * b[i];  
    }  
    return sum;  
}
```

# Protect shared data

The “critical” pragma allows only one thread at a time to execute the update of sum.

```
float dot_prod(float* a, float* b, int N)
{
    float sum = 0.0;
#pragma omp parallel for shared(sum)
    for(int i=0; i<N; i++) {
#pragma omp critical
        sum += a[i] * b[i];
    }
    return sum;
}
```

# Reduction Clause

```
#pragma omp parallel for reduction(+:sum)
    for(i=0; i<N; i++) {
        sum += a[i] * b[i];
    }
```

- Local copy of *sum* for each thread
- All local copies of *sum* added together and stored in “global” variable

Operator	Initial Value
+	0
*	1
-	0
^	0

Operator	Initial Value
&	$\sim 0$
	0
&&	1
	0

# Scheduling Clauses

**schedule(static [,chunk])**

**schedule(dynamic[,chunk])**

**schedule(guided[,chunk])**

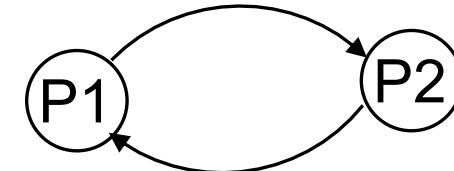
Schedule Clause	When To Use
STATIC	Predictable and similar work per iteration
DYNAMIC	Unpredictable, highly variable work per iteration
GUIDED	Special case of dynamic to reduce scheduling overhead

# Inter-process Communication (IPC)

# IPC: Information sharing



Shared data approach



Message passing approach

# Message Passing System

- A *message-passing system* is a sub-system of a **distributed system** that provides a set of **message-based IPC protocols** and does so by shielding the details of complex **network protocols** and **multiple heterogeneous platforms** from programmers.

# A good message passing system

- Simplicity
  - Uniform semantics
  - Efficiency
  - Reliability
  - Correctness
  - Flexibility
  - Security
  - Portability
- A typical message structure
- Header
    - Addresses
      - ✓ Sender address
      - ✓ Receiver address
    - Sequence number
    - Structural information
      - ✓ Type
      - ✓ Number of bytes
  - Message

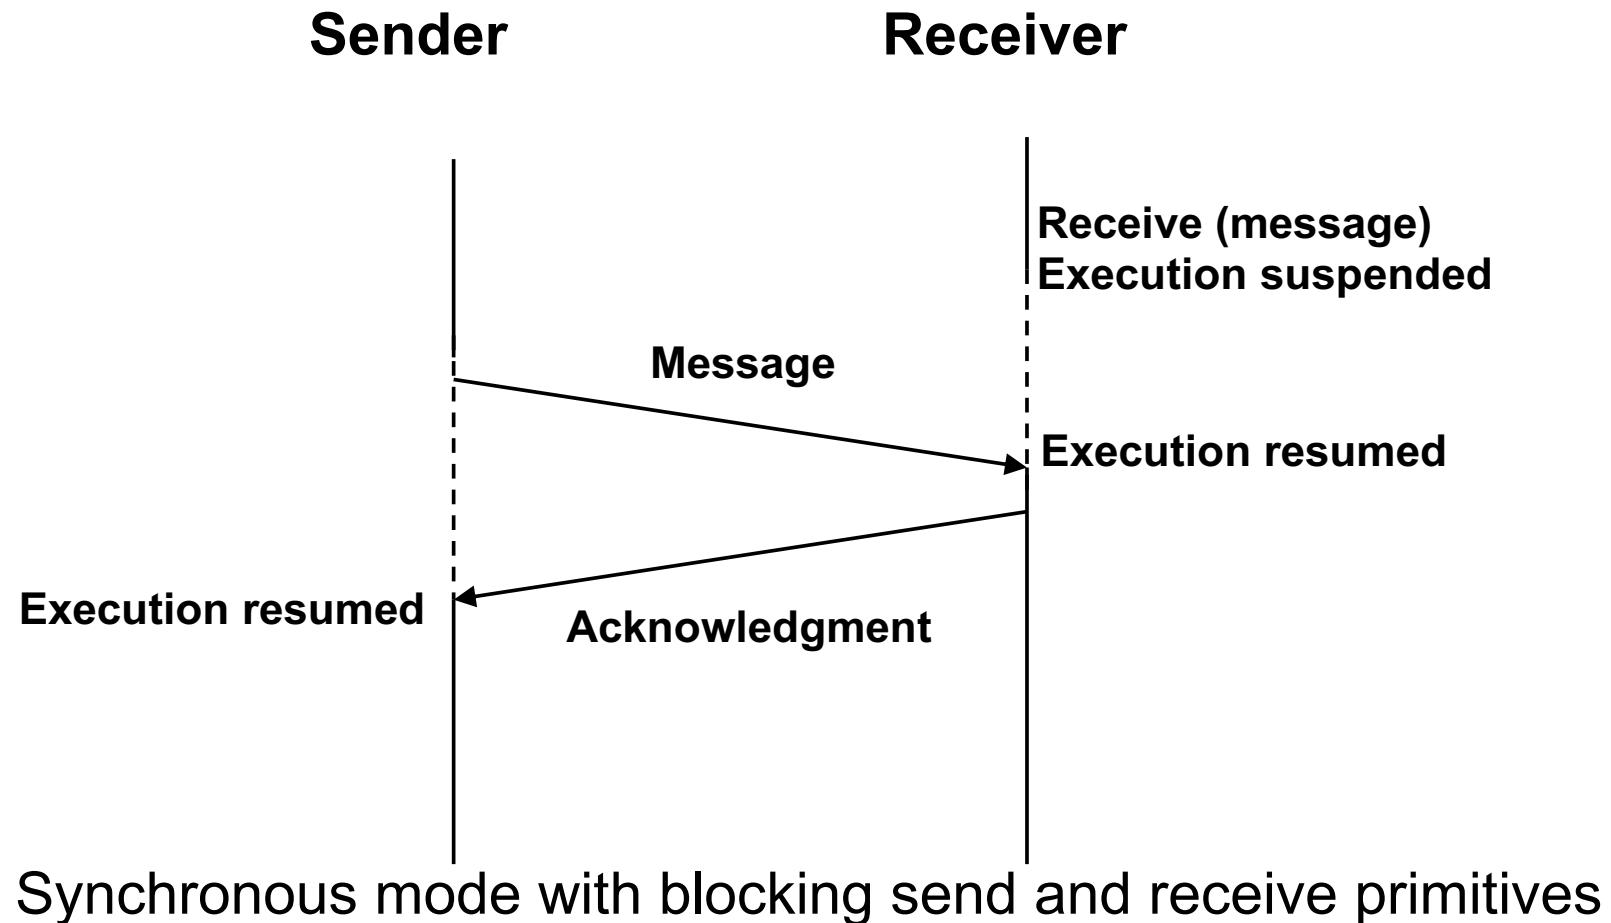
# Issues

- Identity related
- Network Topology related
- Flow control related
- Error control and channel management

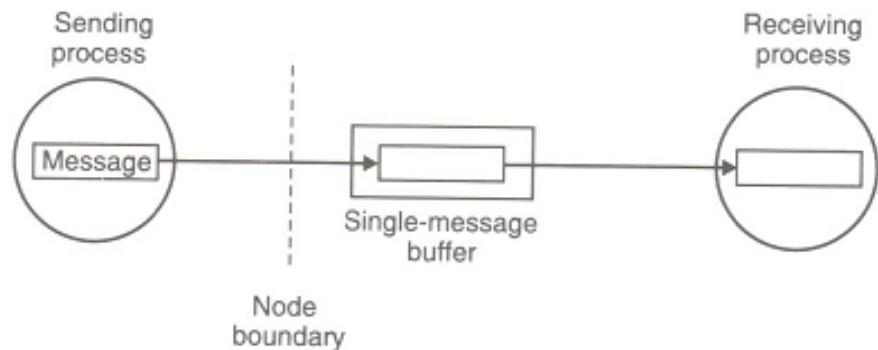
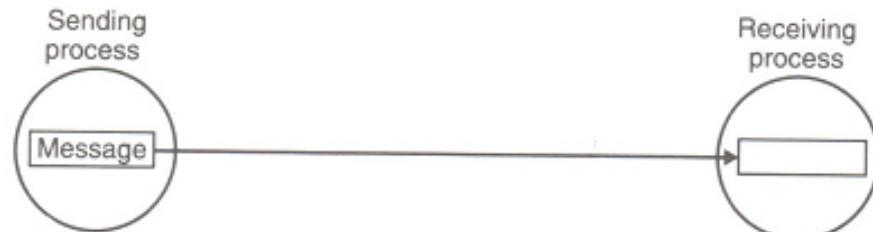
# Synchronization in IPC

- Send primitive
  - Blocking
  - Non blocking
- Receive primitive
  - Blocking
  - Non Blocking
    - Polling
    - interrupt

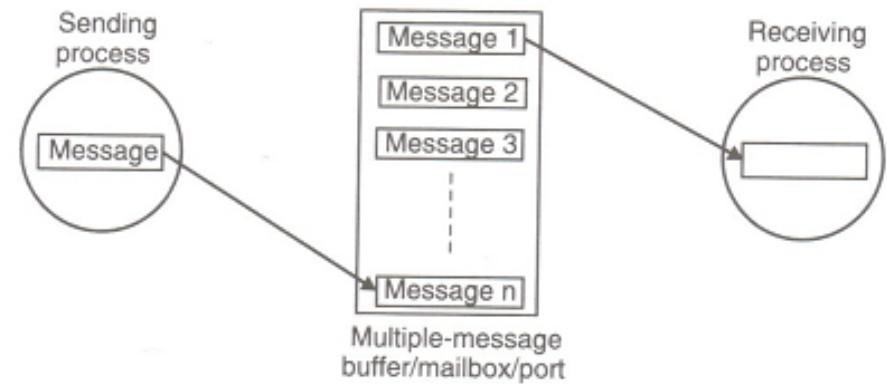
# Synchronous & Asynchronous Communication



# Buffering



Synchronous System

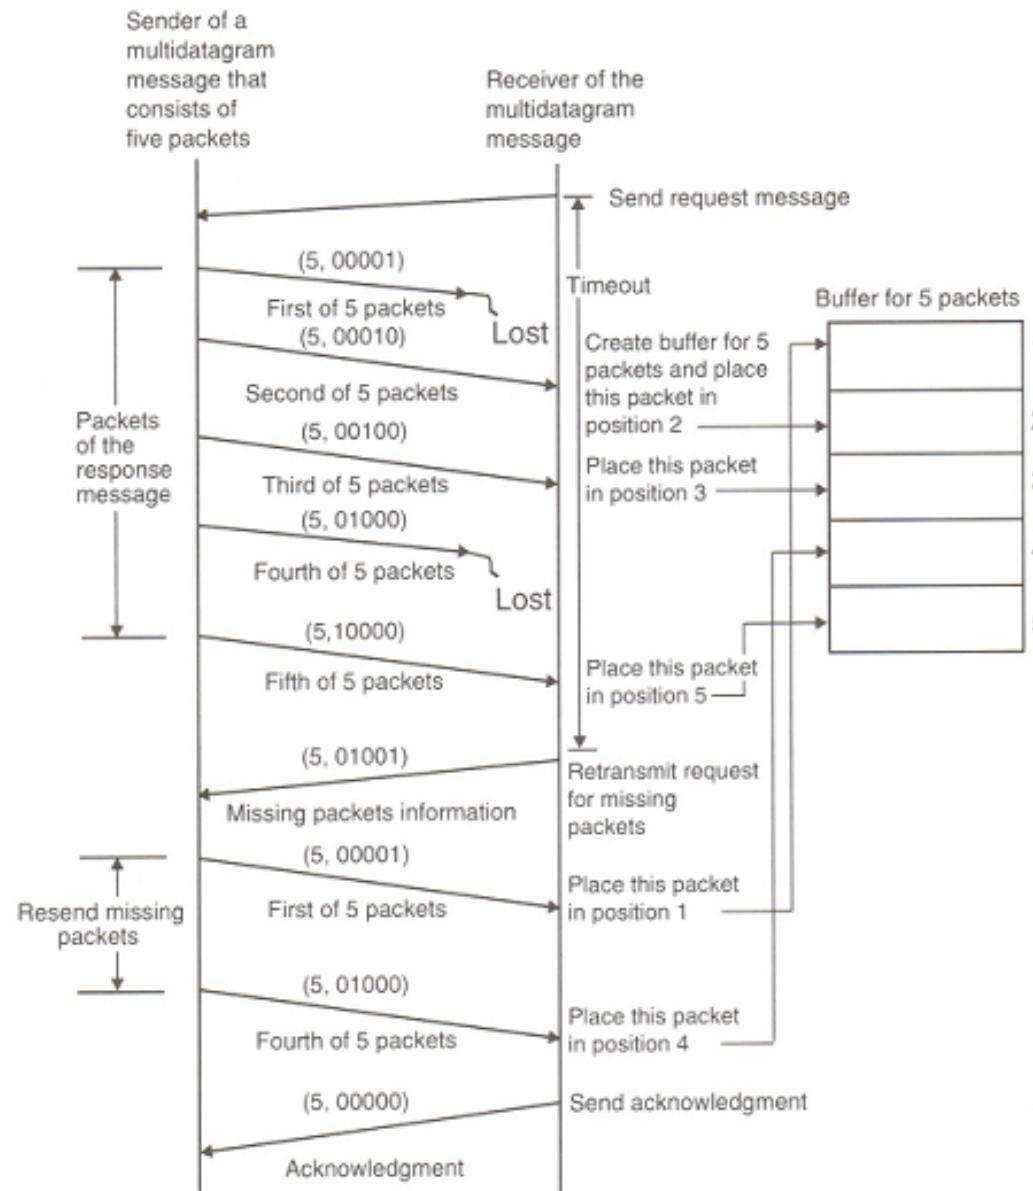


Asynchronous System

# Multi-datagram Messages

- MTU (maximum transfer unit)
- Messages may be single-datagram messages or multi-datagram messages
- Assembling and disassembling is the responsibility of message passing system.

# Using Bitmap for multidatagrams

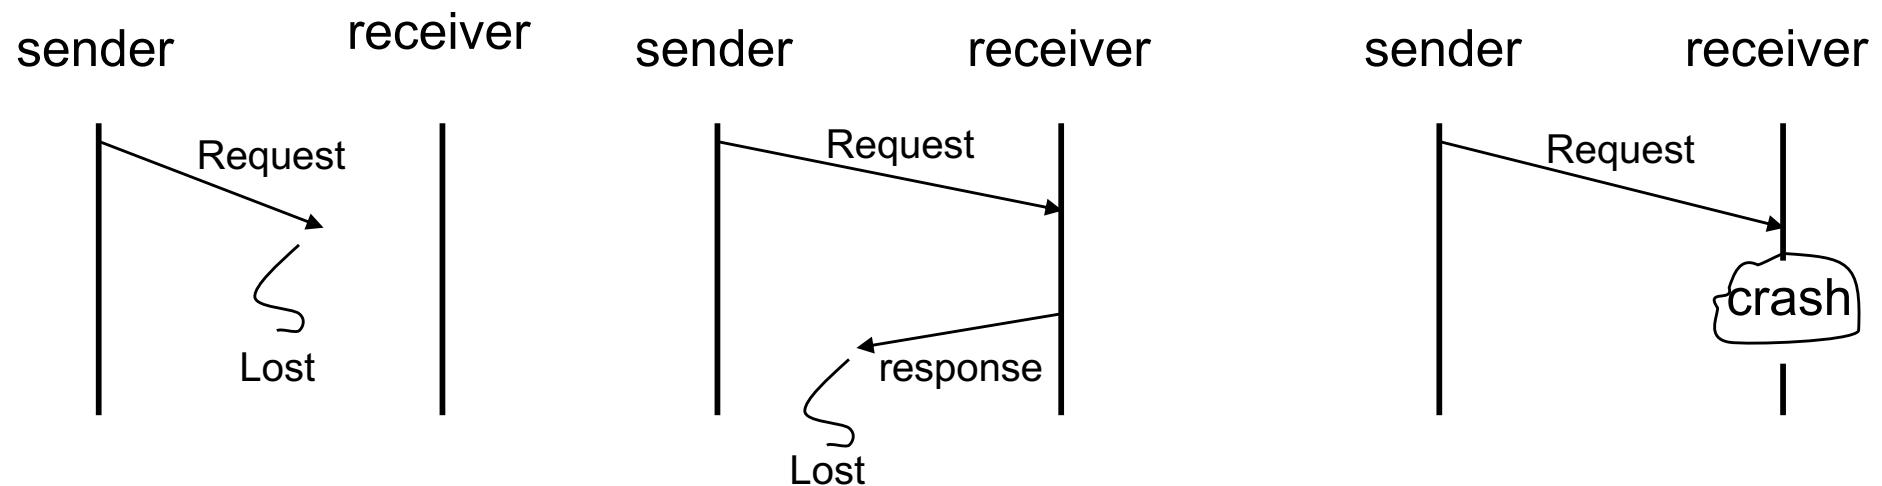


# Encoding Decoding

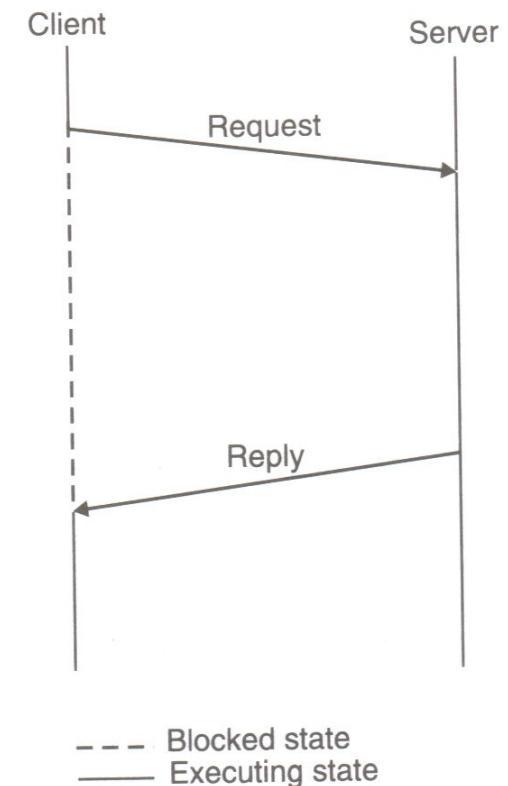
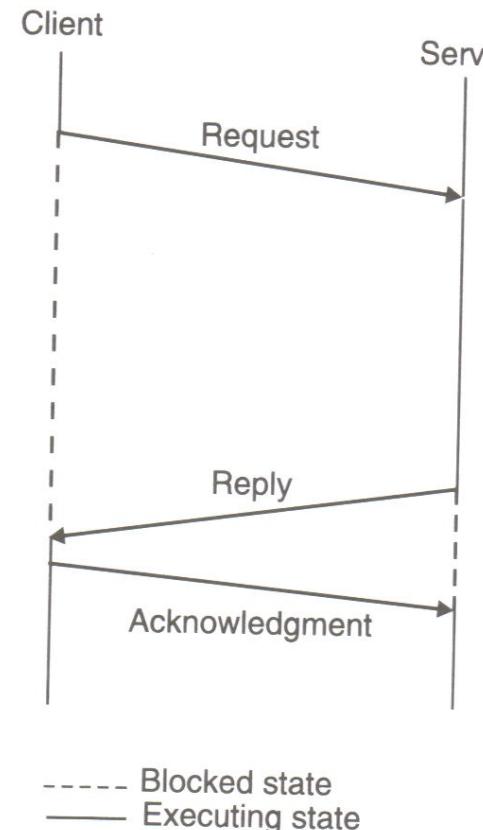
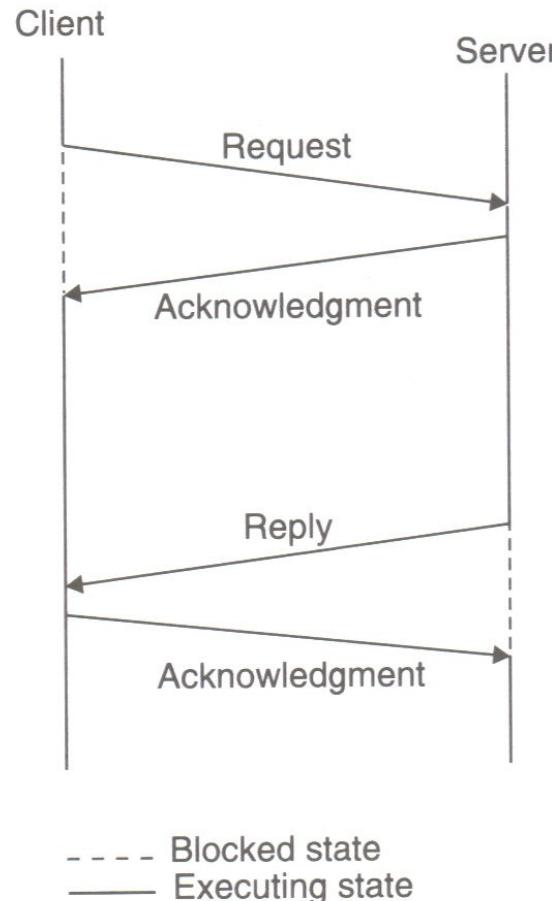
- Encoding/Decoding is needed if sender and receiver have different architecture
- Even for Homogeneous Encoding/Decoding is needed for using an absolute pointer and to know which object is stored in where and how much storage it requires

# Failure handling

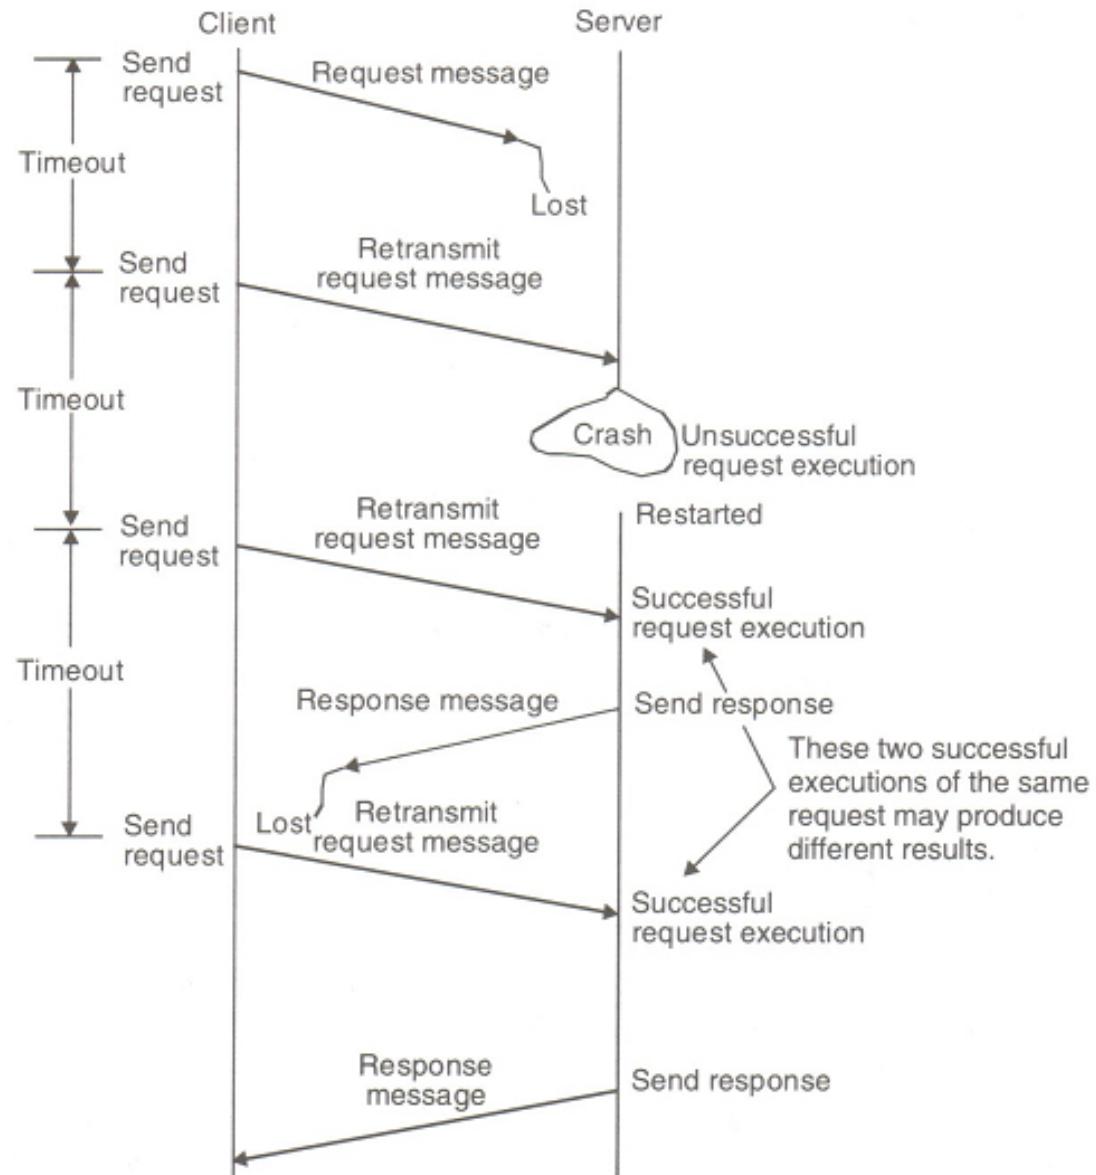
- Failure classification
  - Loss of request message
  - Loss of response message
  - Unsuccessful execution of the request



# Blocking



# Fault Tolerant System



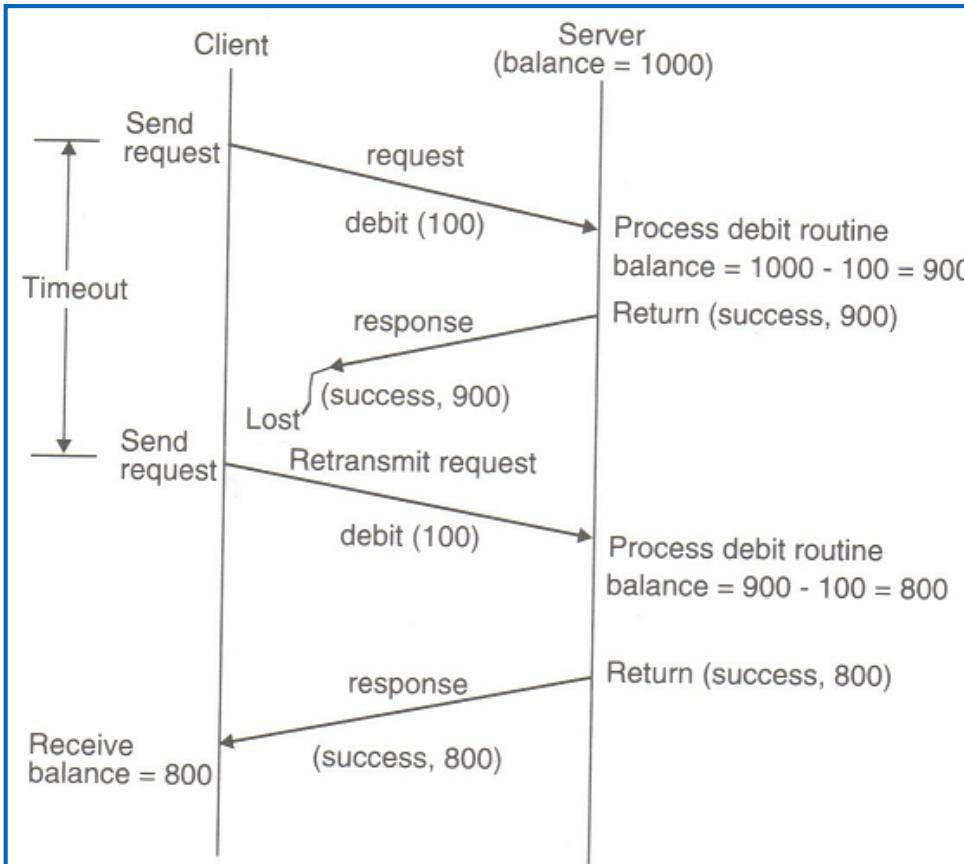
# Idempotency

- What is the difference between the following two functions?

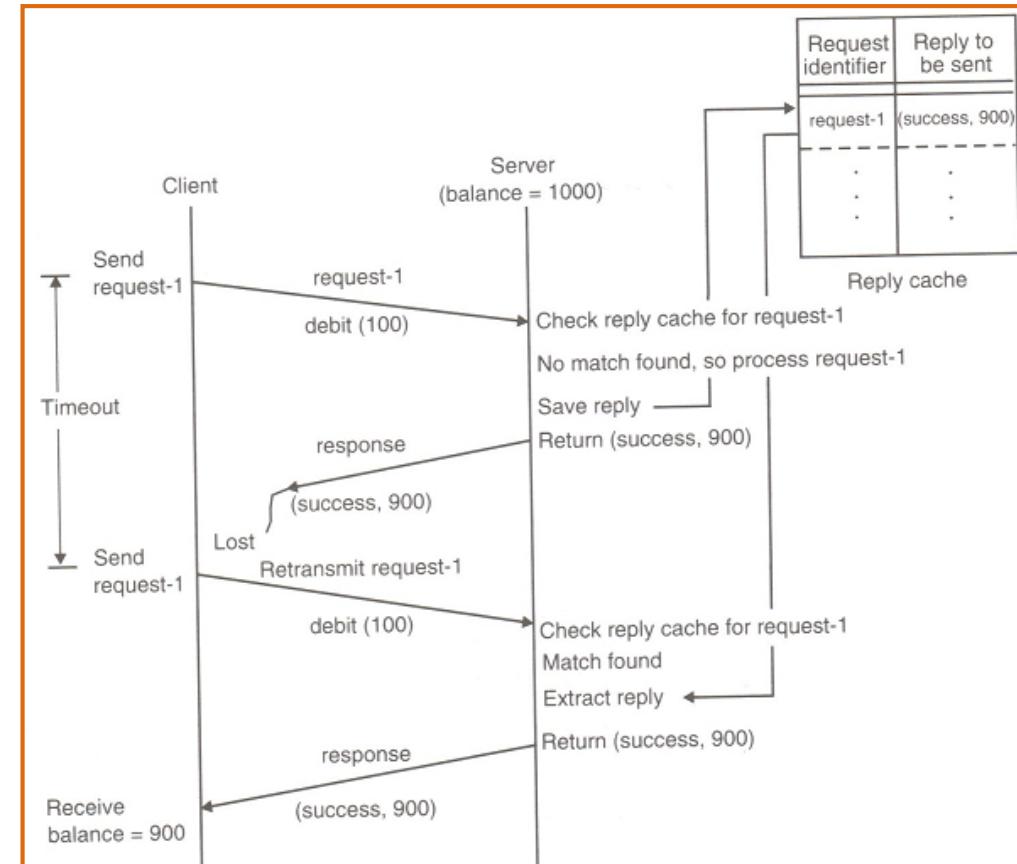
```
getSqrt(n)
{ return sqrt(n)
}
```

```
Debit (amount)
{ if (balance>amount)
  balance= balance-amount
  return ("success", balance)
  else return ("failure", balance)
}
```

# Implementation of Idempotency



–Non-Idempotent



- Adding sequence number with the request message
- Introduction of ‘Reply cache’

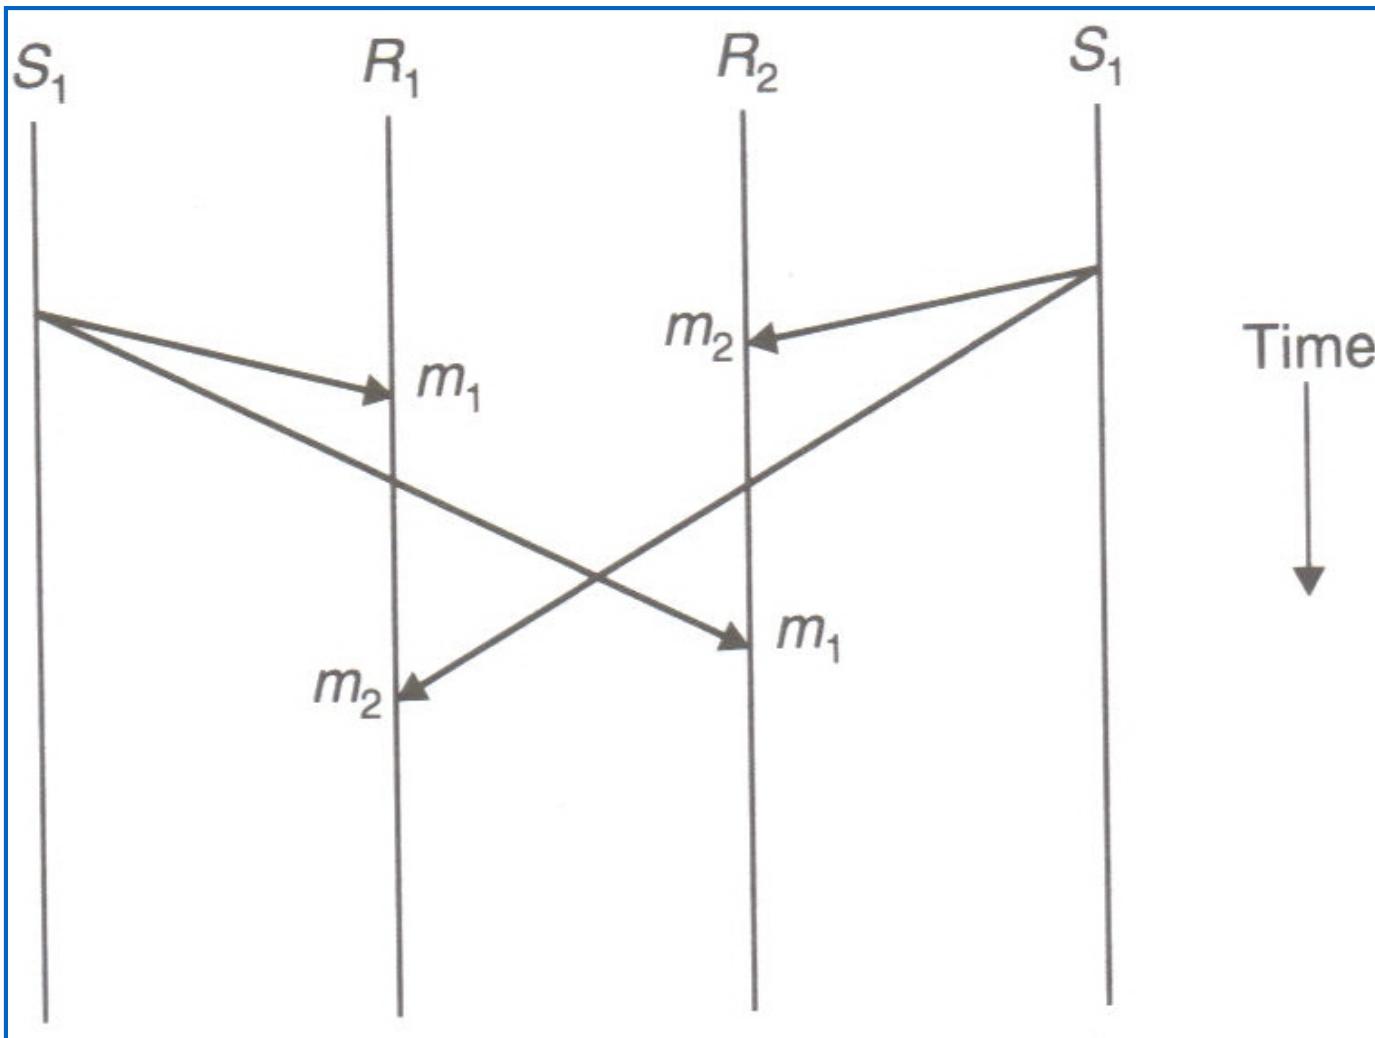
# Group Communication

- Group communication:
  - One to many, Many to one, Many to many
- One to many
  - Group management, Group addressing
  - Buffered and unbuffered multicast
  - Send-to-all and Bulletin-Board semantics
  - Flexible reliability in multicast communication
  - Atomic multicast

# Many to Many communication

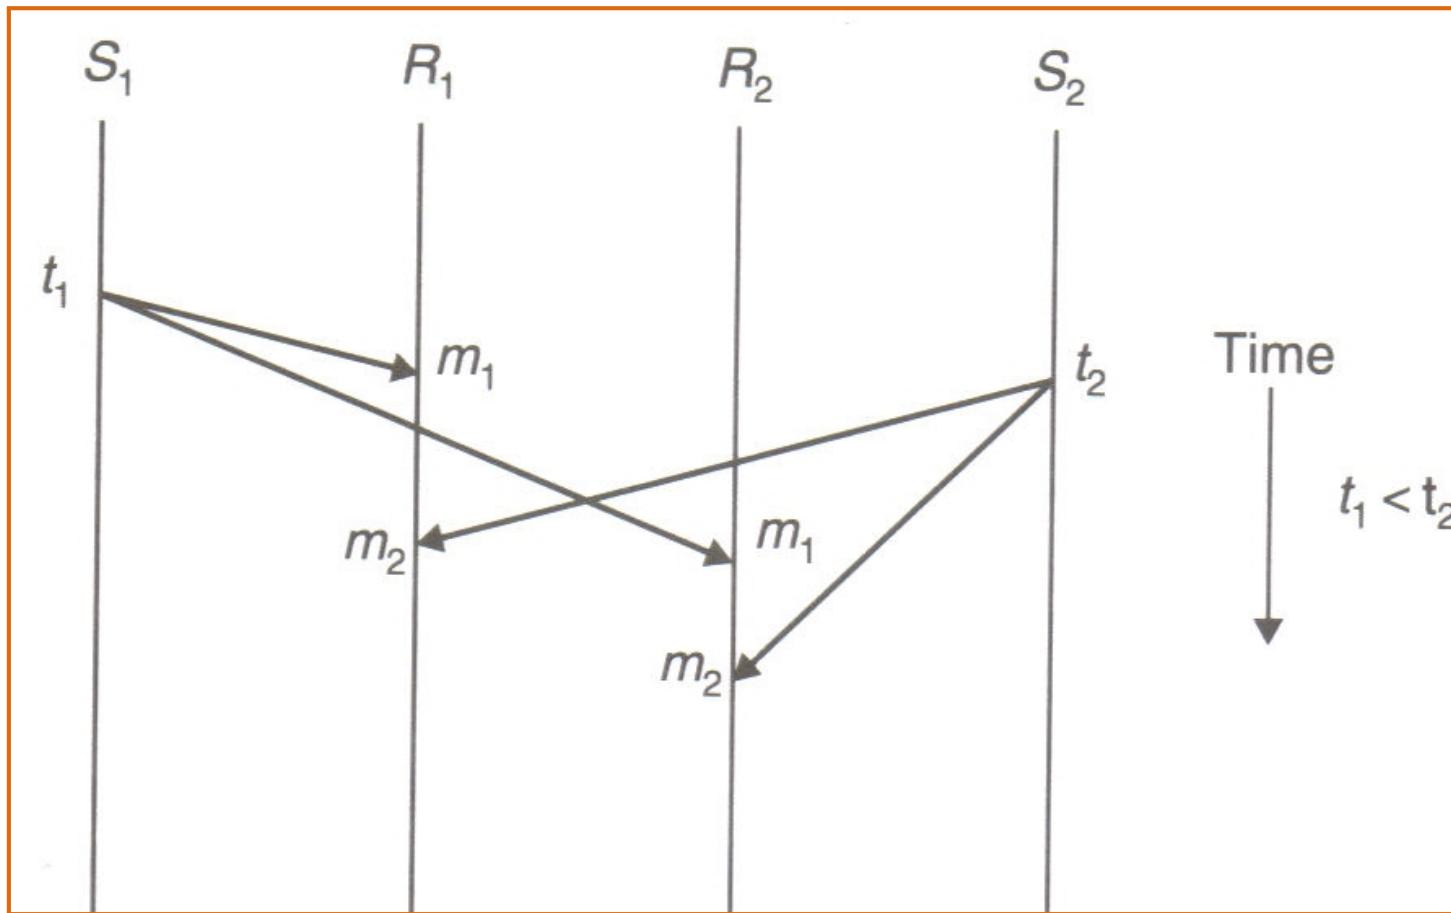
- *Ordered message delivery* is an important issue.
- For example, two server processes are maintaining a single salary database. Two client processes send updates for a salary record. What happen if they reach in different order? (will sequencing of messages help in this case?)
-

# Ordering Semantics



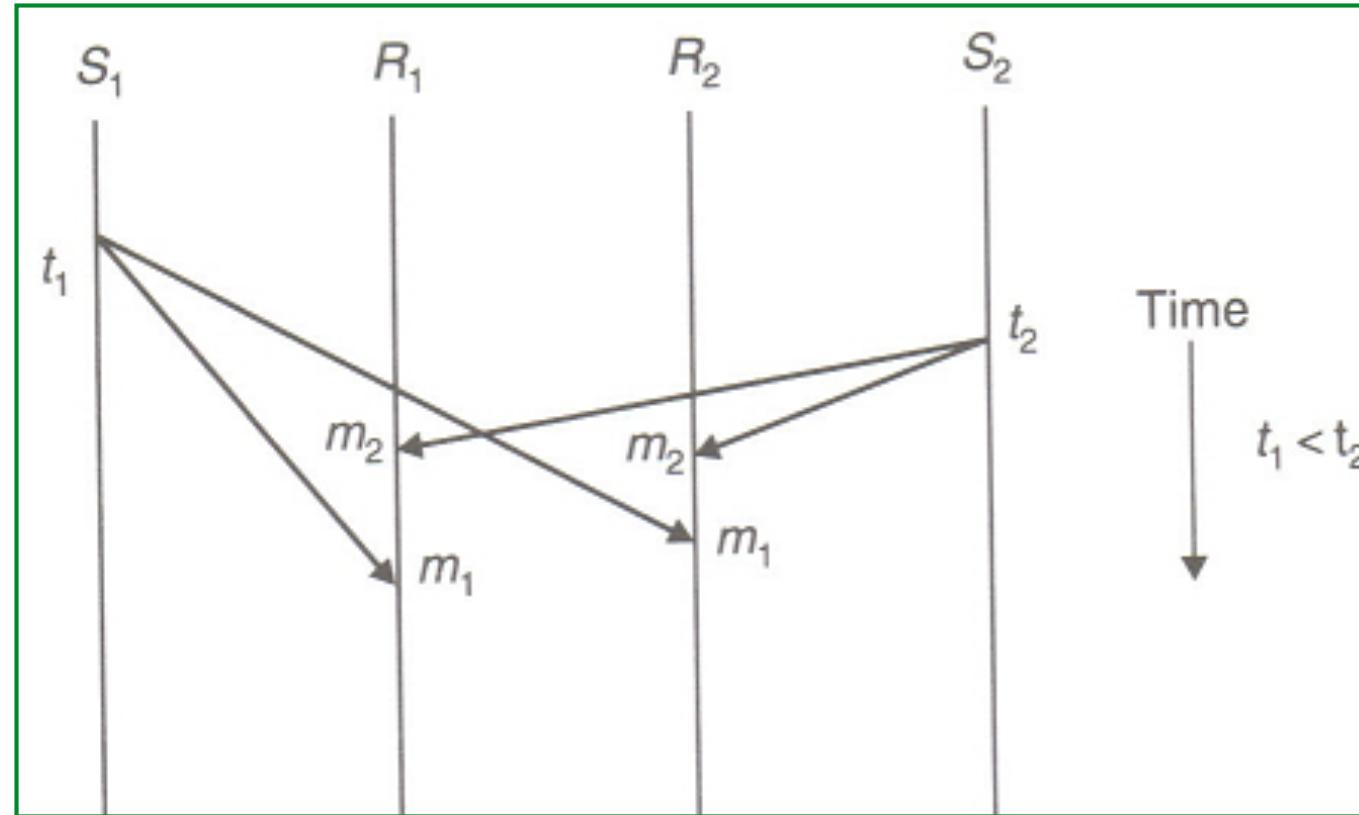
No ordering constraint

# Absolute Ordering Semantics



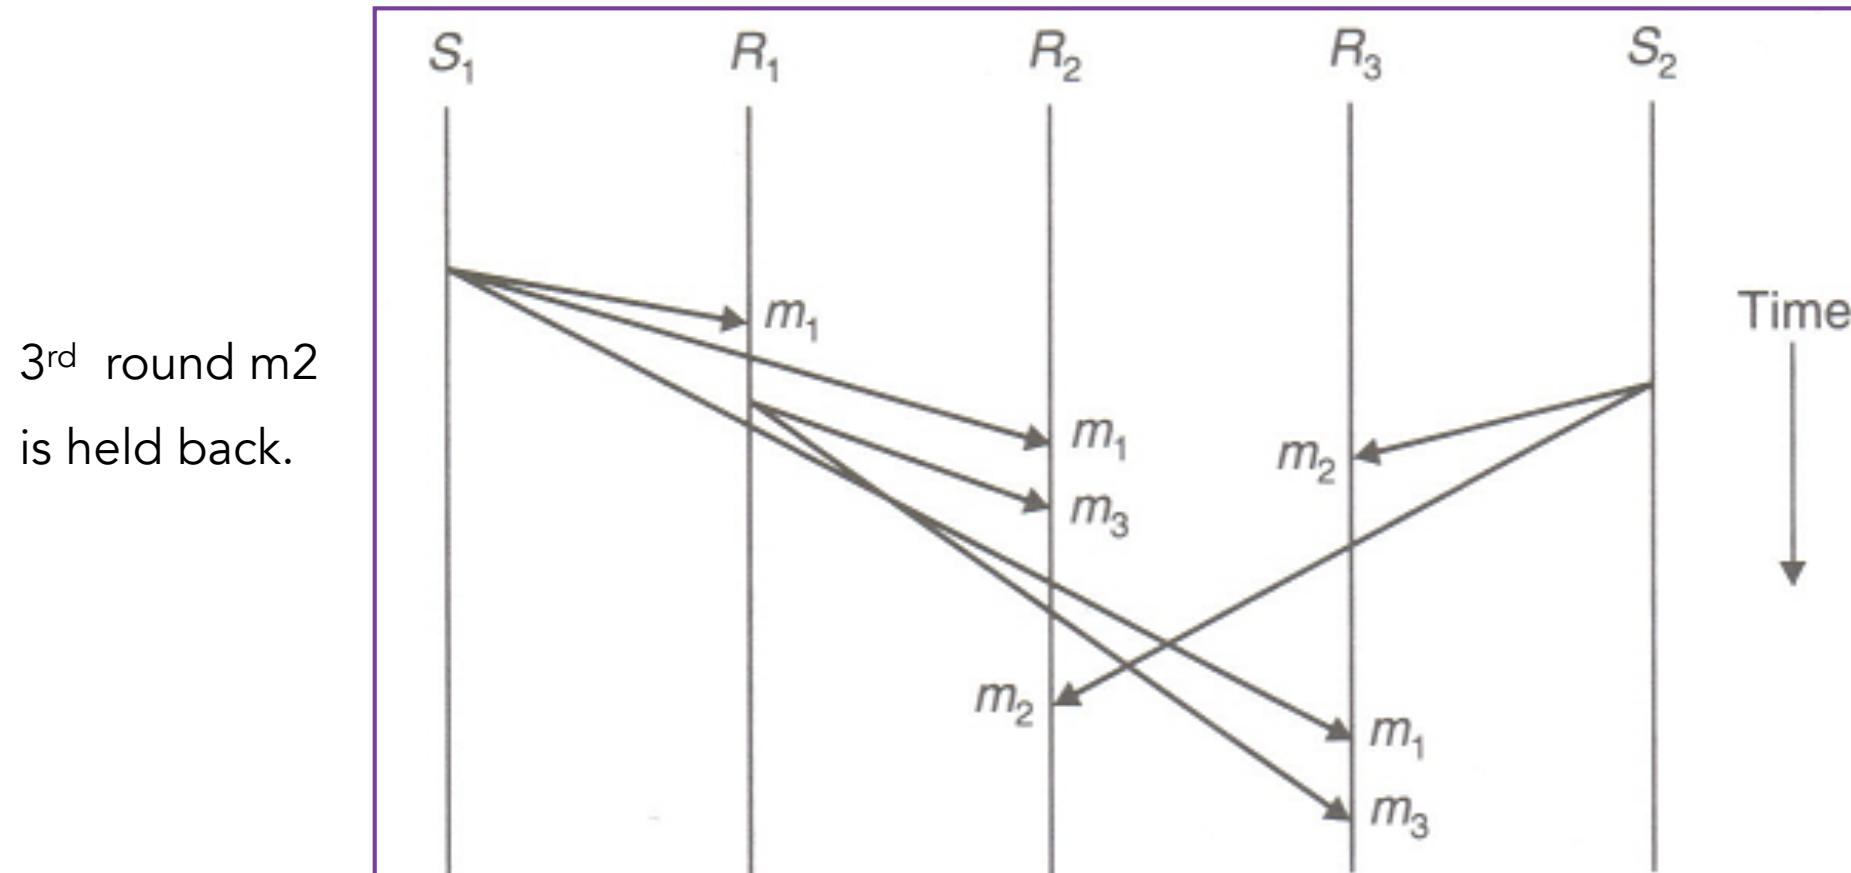
All messages are delivered to all receiver processes in the **exact order** in which they were sent. Using global timestamp as message identifiers with **sliding window protocol**

# Consistent Ordering Semantics



All messages are delivered to all receivers in the same order. However, this order may be different from the order in which messages were sent.

# Causal Ordering Semantics



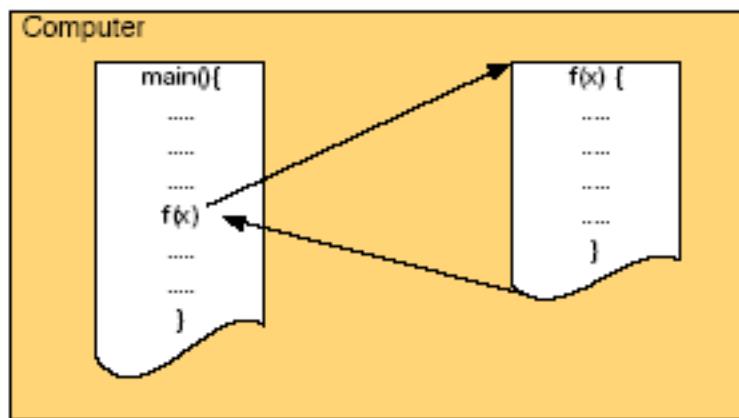
If the event of sending one message is **causally related** to the event of sending another message, the two messages are delivered to all receivers in correct order. Two message sending events are said to be causally related if they are correlated by the **happened-before relation**.

# Remote Procedure Call (RPC)

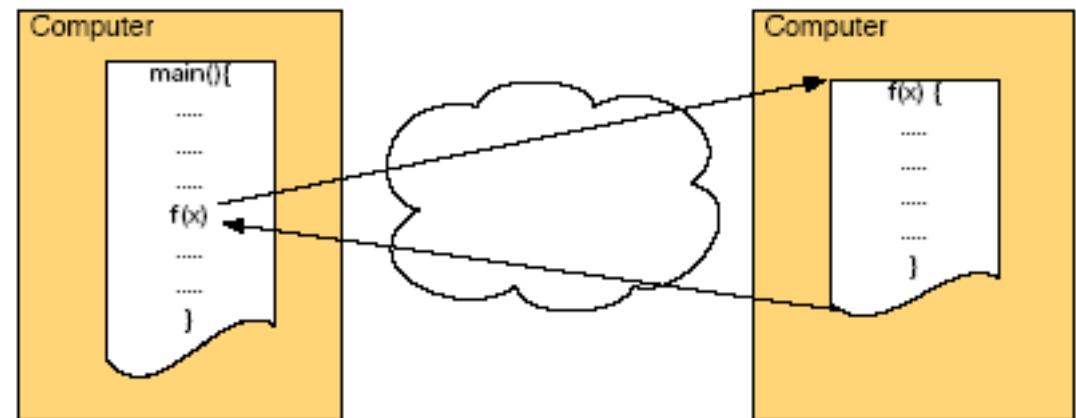
# Features of RPC

- Simple call syntax
- Familiar semantics
- Specification of a well defined interface
- Ease of use
- Generality
- Efficiency

# Local and Remote Procedure Call

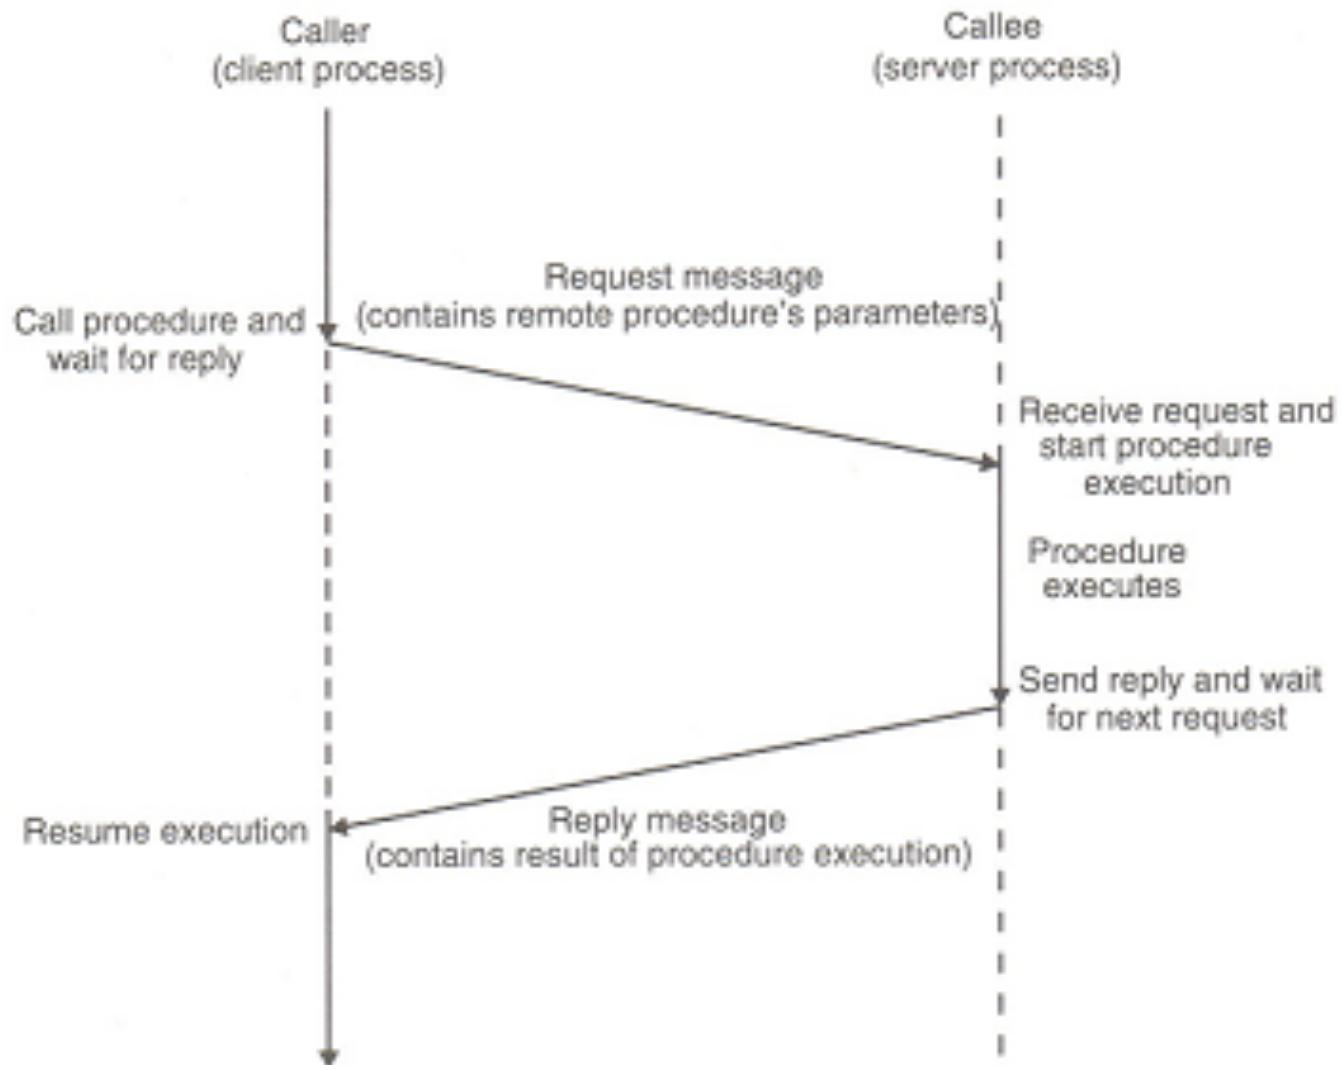


Local Procedure Call



Remote Procedure Call

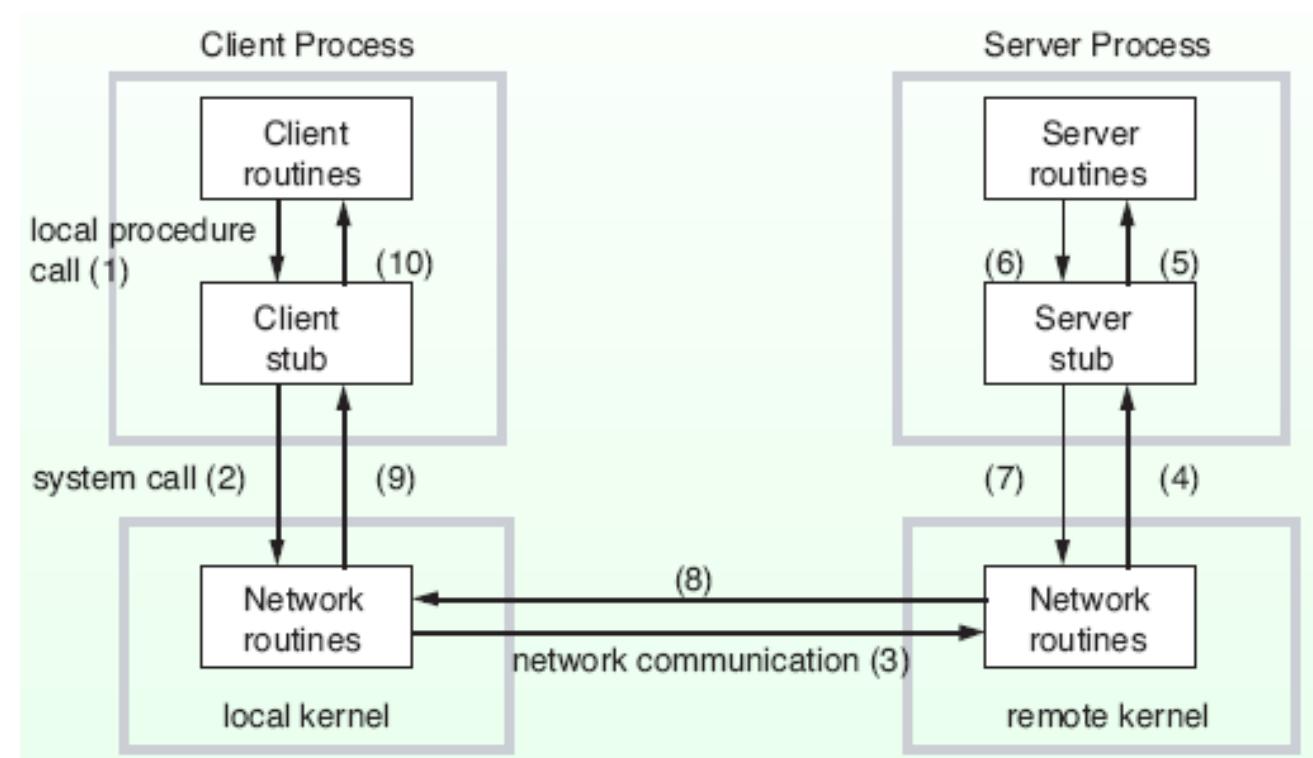
# A Typical Model of RPC



# Implementing RPC mechanism

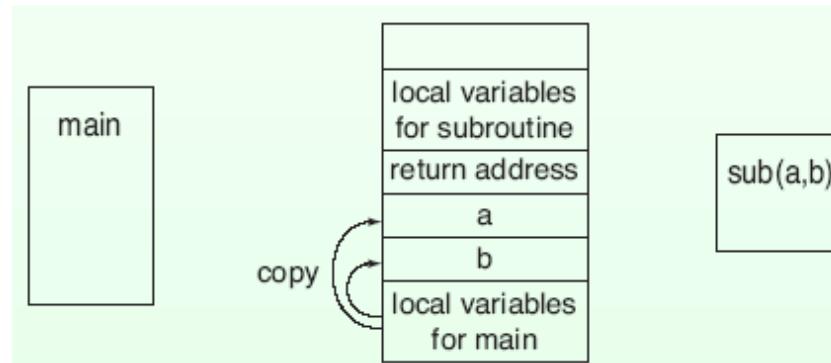
RPC implementation involves five elements:

- The client
- The client stub
- The RPCRuntime
- The server stub
- The server

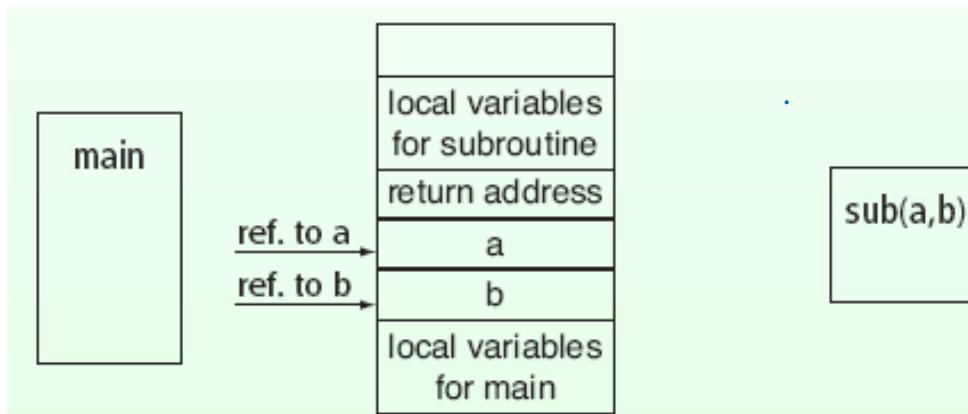


# Parameter Passing Mechanisms

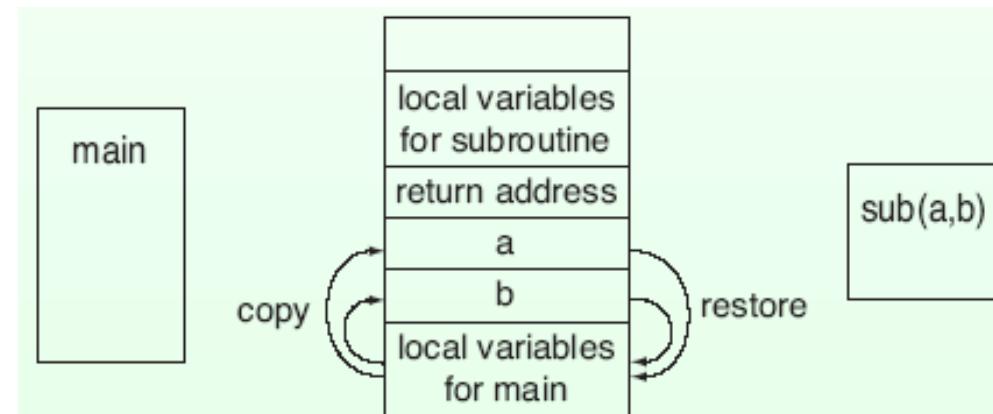
- When a procedure is called, parameters are passed to the procedure as the arguments. There are three methods to pass the parameters.



• call-by-value

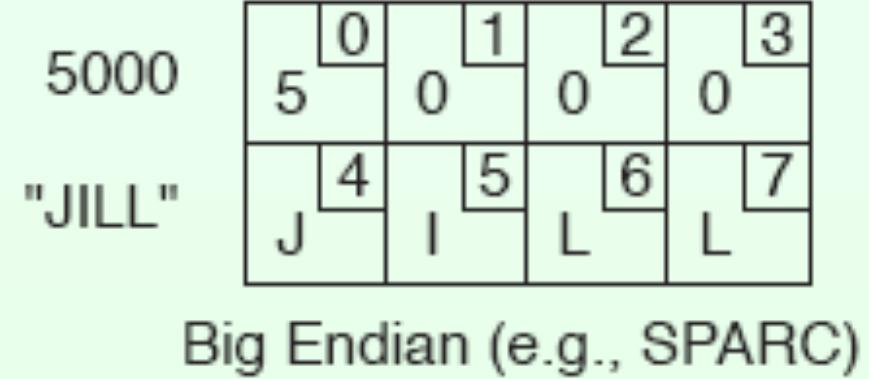
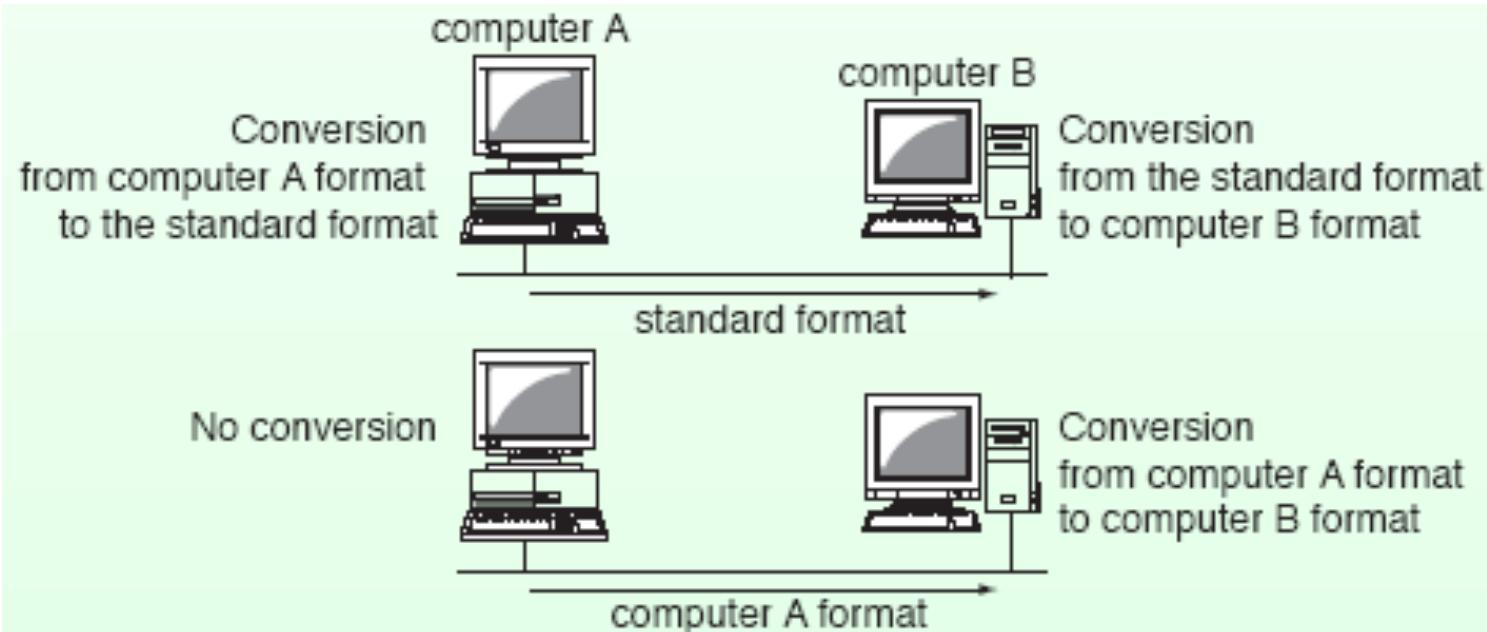


• call-by-reference

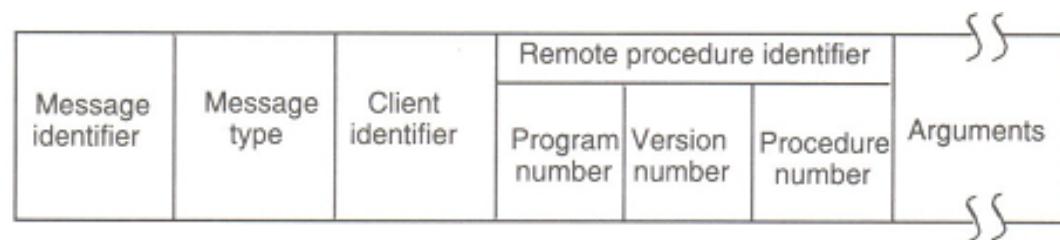


• call-by-copy/restore

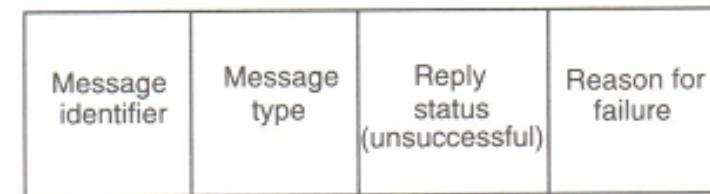
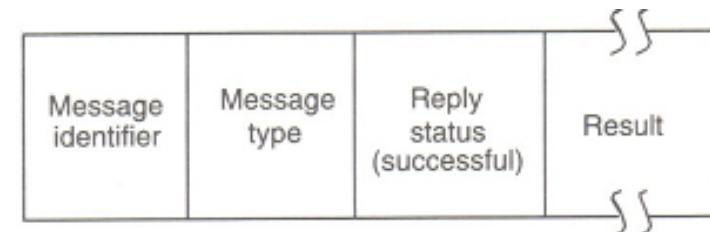
# Transmit over the network



# RPC Messages

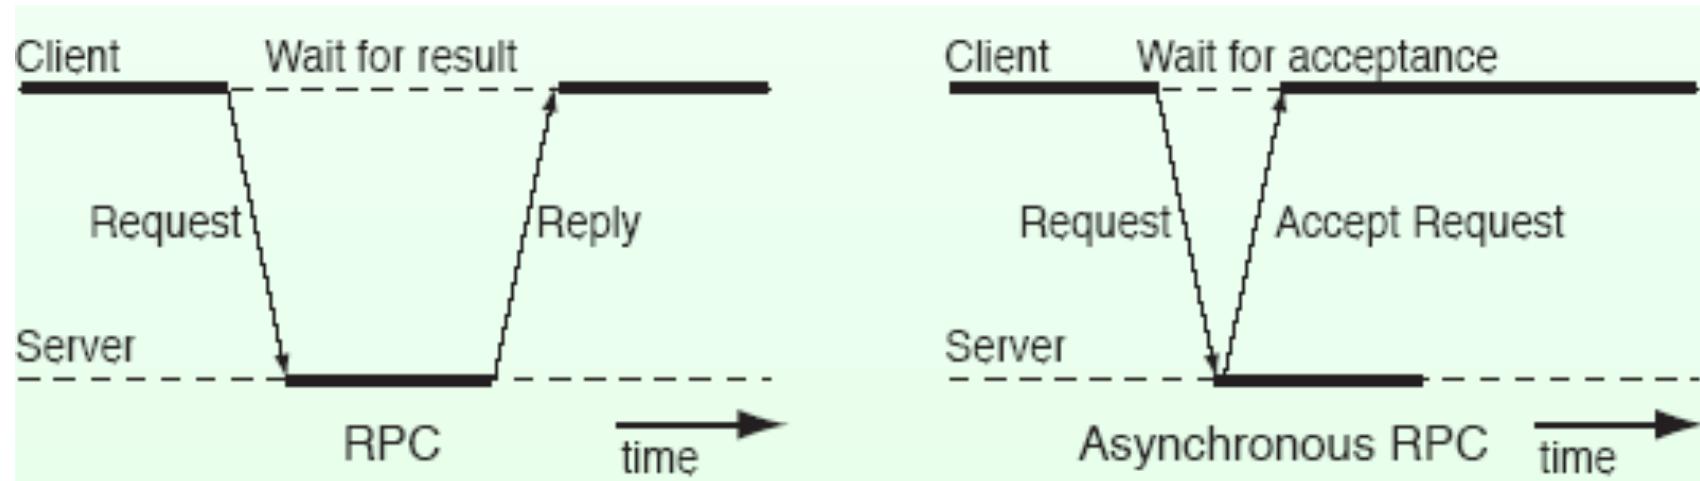


A typical RPC Call message format

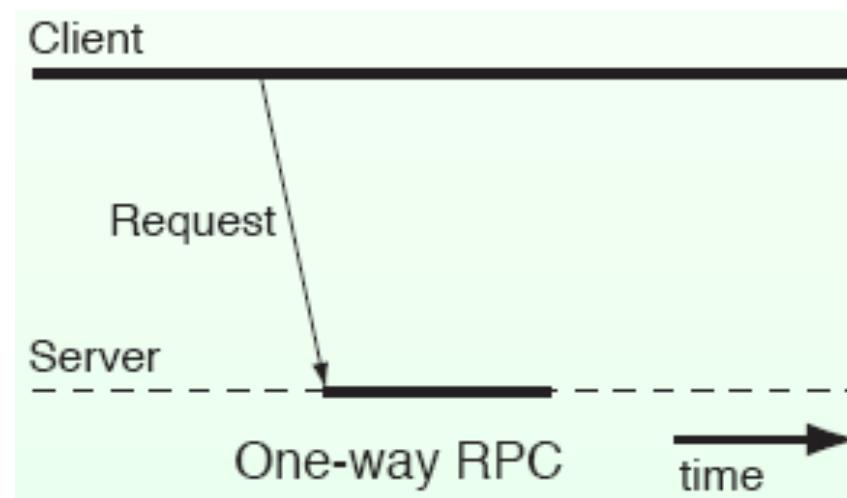


A typical RPC Reply message format  
(successful and unsuccessful)

# Variations of RPC



## Asynchronous RPC



## One-way RPC

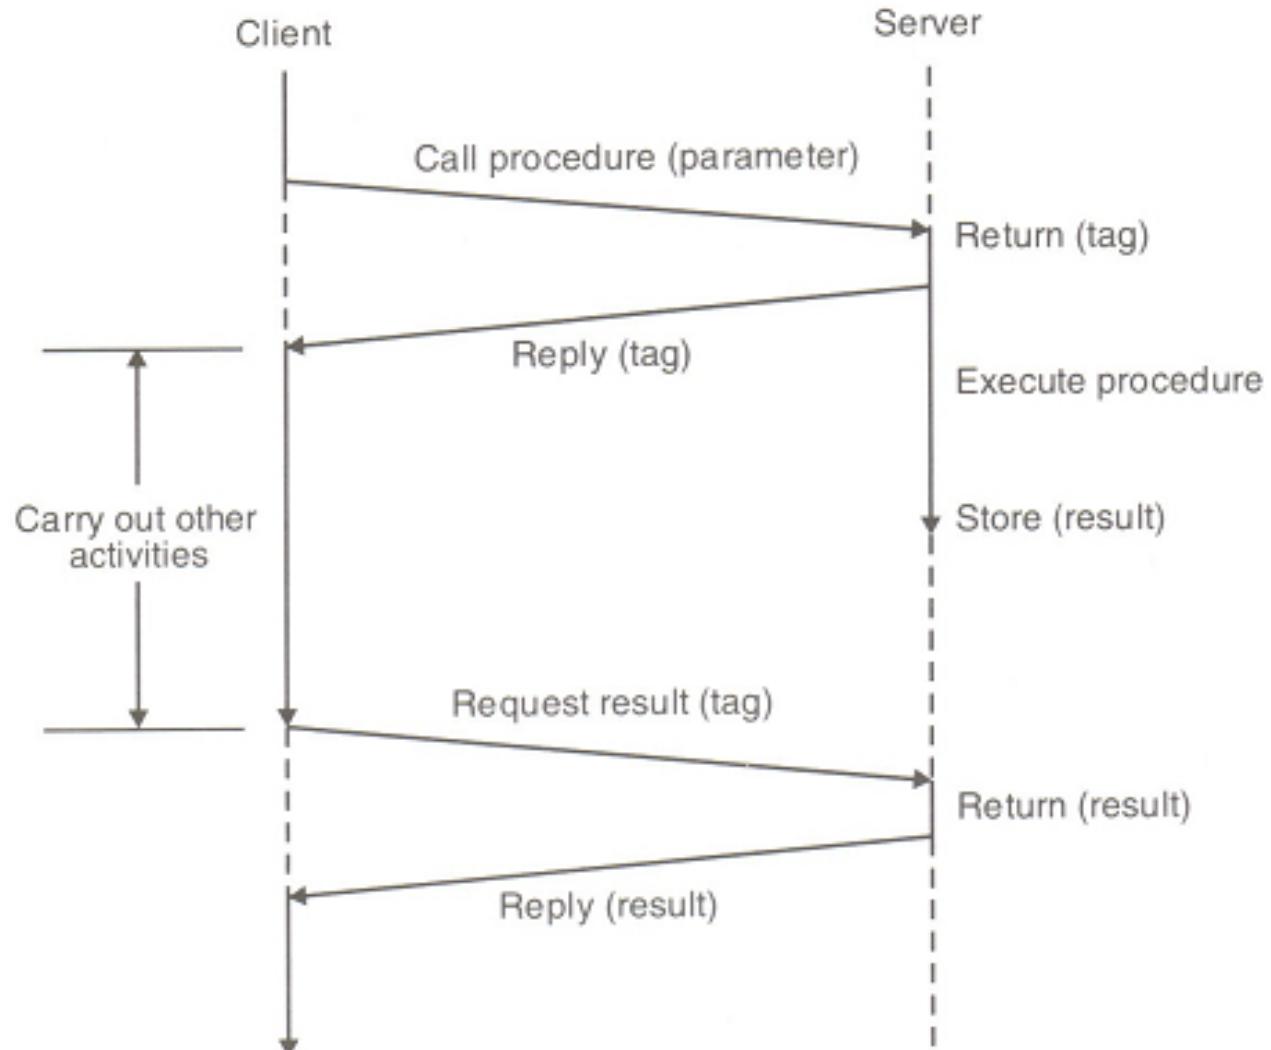
# Optimizations for better Performance

- Six different ways:
  - Concurrent access to multiple servers
  - Serving multiple requests simultaneously
  - Reducing per-call workload of servers
  - Reply caching of idempotent remote procedures
  - Proper selection of timeout values
  - Proper design of RPC protocol specification

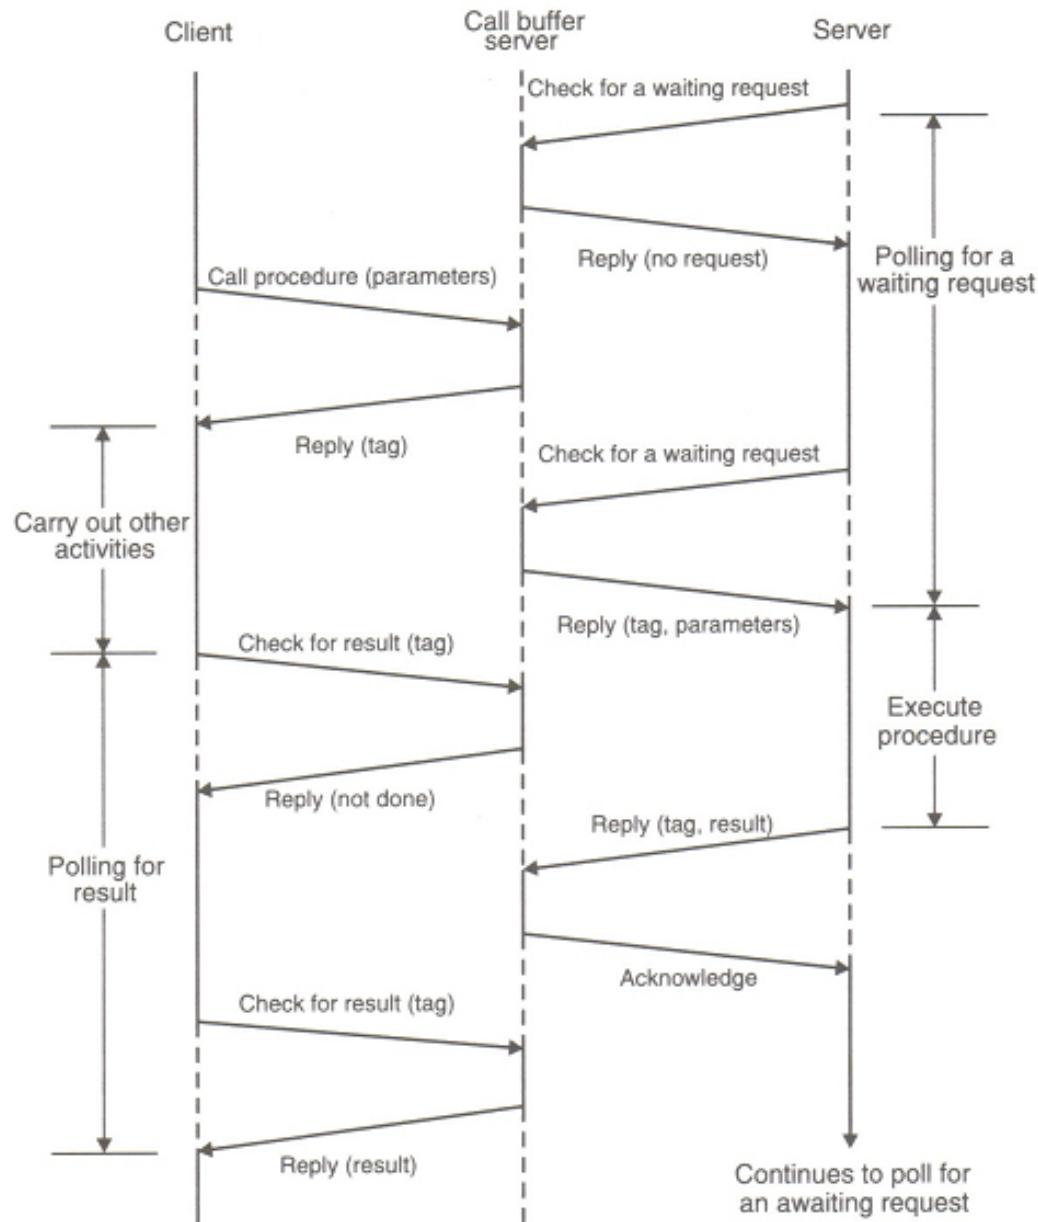
# Concurrent Access to Multiple Servers

- One of the following three may be adopted:
  - Threads
  - Early reply approach [Wilbur and Bacarisse]
  - Call buffering approach [Gimson]

# Early Reply Approach



# Call Buffering Approach



# Serving Multiple Requests Simultaneously

- Following types of delays are common:
  - A server, during the course of a call execution, may wait for a shared resource
  - A server calls a remote function that involves computation or transmission delays
- So the server may accept and process other requests while waiting to complete a request.
- **Multiple-threaded server** may be a solution.

# FLUX: Multithreading

**Multithreading**

**JYGAFY**

# References

- Birman, K. P. and Renesse, R. V. Reliable Distributed Computing with the ISIS Toolkit. IEEE Computer Society Press, 1994.
- Birrell, A. D. and Nelson, B. J. Implementing remote procedure calls. ACM Trans. Comput. Syst. 2(1), 39-59, 1984.
- Wilbur, S. and Bacarisse, B. *Building distributed systems with remote call*, Software Engineering Journal, 2(5), 148-159, 1987.
- R. Gimson. Call buffering service. Technical Report 19, Programming Research Group, Oxford University, Oxford University, Oxford, England, 1985.
-

# Lab Week 3 Overview

- Write a serial program
- Convert to parallel program using POSIX Threads

# Tutorial Week 4 Overview

- Inter-Process Communications
- RPC

# Next week: Distributed Memory

- Parallel Computing in Distributed Memory
- Message Passing Library