

MonARCH Welcomes Students from the FIT3143 Parallel Computing Course

August 2020

Philip Chan
Monash HPC Team



Learning Outcomes

1. Better understanding of HPC

- a. how to access MonARCH system
- b. what do you need installed to access a compute cluster
- c. account information, initial password, and signing into MonARCH

2. Basic Linux and Job Submission

- a. navigation, file management, directory management
- b. \$HOME and Project Folders on MonARCH
- c. running the text editor
- d. managing data transfers

3. How to use a compute cluster

- a. basic concepts
- b. SLURM commands
- c. submitting your first job, checking results

4. Further reading

Introduction to High-Performance Computing

- High-performance computing is about employing the best and **cost-effective** technologies, from processors, storage and networks, to provide **aggregated** computational capabilities beyond what is typically available to the end-user
- HPC versus HTC
 - **high-performance** -- running a single program as quickly as possible
 - **high-throughput** -- running many independent programs as quickly as possible
- HPC/HTC is predominantly achieved by one or more **compute clusters**
- HPC/HTC are enabling technologies for larger experiments, more complex data analyses, achieving higher accuracy in computational models

Introducing MonARCH

MonARCH (Monash Advanced Research Computing Hybrid)

- university cluster for high-performance/high-throughput computing
- open for access to Monash researchers and collaborators
- hosts courses in Chemistry, Engineering and FIT that require access to advanced tools or hardware, like lots of CPU cores, or GPU devices
- made up of over 83 servers connected via a fast network
- runs on top of the Monash internal cloud
- over 233 active users in the past 90 days
- 40-50 active users who together run 400-500 calculations at once

What you need to access MonARCH

On Windows you will need two applications:

On Linux and Mac OS systems, the commands `ssh` and `scp` are pre-installed
FileZilla is available on these platforms

- an **ssh** client
 - putty
 - <https://the.earth.li/~sgtatham/putty/latest/w64/putty.exe>
 - MobaXterm
 - https://download.mobatek.net/2032020060430358/MobaXterm_Portable_v20.3.zip
- a secure copy client (for data file transfers)
 - WinSCP
 - <https://winscp.net/download/WinSCP-5.17.7-Portable.zip>
 - FileZilla (choose the appropriate OS)
 - https://filezilla-project.org/download.php?show_all=1

Where to get info about my MonARCH Account

FIT3143 staff and students are assigned individual user accounts on MonARCH; and your user account is **distinct** from your Monash username.

Your Monash username is the word **before** the @ in your email address, e.g.:

abcd0001@student.monash.edu

Use **+** on MobaXterm => new connection => enter this as the **address**:

118.138.234.199

When prompted for a username, enter your Monash username; then enter your **MyMonash** password. Read the message displayed.

Logging into the MonARCH cluster

Keep your password secure, do not share this with anyone.

By now, you should know your assigned MonARCH (HPC) username:

it is prefixed by “fit” followed by three decimal digits, e.g., fit015

On MobaXterm, create a new connection using this address:

monarch.erc.monash.edu

with your assigned HPC username (e.g., fit015) and then enter your given password.

On successful SSH access into MonARCH

You will see a “Welcome to MonARCH” MOTD (message-of-the-day)

If there are any **updates** on cluster availability, we will be posting the information on the MOTD.

Below the MOTD, you should see a command prompt, e.g.,

```
[fit015@monarch-login1 ~]$
```

awaiting your textual commands.

```
*****
* Welcome to MonARCH                                     *
*                                                         *
* For assistance, please contact: mcc-help@monash.edu    *
* Documentation: https://docs.monarch.erc.monash.edu     *
*****
```

```
- Useful SLURM Commands:
  squeue
  SQ                               for extended info on jobs
  sbatch <slurm_script_file>
  scontrol show job <JOBID>
  scancel <JOBID>

- M3/MonARCH User Scripts:
  show_job
  show_job <JOBID>
  show_cluster
```

```
*****
Please limit CPU-intensive processes on this login node
to a maximum of 20 minutes. Please use the sbatch command
to submit longer processes to the queue. There is a short
partition for jobs that take a maximum of 24 hours.
*****
```

```
NOTE: Your respective $HOME folders are currently backed-up.
The Lustre project & scratch folders are NOT backed-up.
Please ensure that you keep a copy of IMPORTANT data
outside of MonARCH. We will soon support Lustre backups.
```

```
*****
```


First Commands on MonARCH

You are now within your “home folder”, we call **\$HOME**, also called **~**

First commands:

- pwd** - displays the working directory (where you are)
- echo \$HOME** - try this command and compare its output with above
- date** - display the system date and time
- cal 2019** - show the calendar for 2019
- fortune** - try this
- clear** - clears the terminal window
- passwd** - change your MonARCH password (do this soon)

Introducing MonARCH

`show_cluster`

- “CPU-only” nodes (three configurations)
- “GPU” nodes (two configurations)
- 1,476 physical CPU cores
- over 400 TB of Lustre storage
- compute nodes are virtual machines on the Monash Research Cloud

Login node: `monarch.erc.monash.edu`

`monarch-login1.erc.monash.edu`

`monarch-login2.erc.monash.edu`

Data Transfer: `monarch-dtn.erc.monash.edu`

Docs: <https://docs.monarch.erc.monash.edu>

gf01	K80	gpu	18	206	3	Running
gp01	P100	gpu	0	153	0	Busy
gp02	P100	gpu	0	142	0	Busy
gp03	P100	gpu	0	166	0	Busy
hs03	CPU	short	0	1	0	Busy
hs04	CPU	short	0	34	0	Busy
hs05	CPU	short	0	34	0	Busy
hs06	CPU	comp	0	54	0	Busy
hs12	CPU	comp	6	63	0	Running
hs13	CPU	comp	0	59	0	Busy
hs14	CPU	comp	0	65	0	Busy
hs16	CPU	comp	0	62	0	Busy
hs17	CPU	comp	0	49	0	Busy
mi00	CPU	comp	4	43	0	Running
mi01	CPU	comp	3	28	0	Running
mi02	CPU	comp	7	10	0	Running
mi05	CPU	comp	4	27	0	Running
mi06	CPU	comp	1	68	0	Running
mi07	CPU	comp	0	11	0	Busy
mi08	CPU	comp	0	26	0	Busy
mi09	CPU	comp	0	21	0	Busy
mi11	CPU	comp	2	7	0	Running
mi12	CPU	comp	6	68	0	Running
mi14	CPU	comp	0	12	0	Busy
mi15	CPU	comp	4	30	0	Running

Summary:

	CPUS		Nodes		K80 GPUs		P100 GPUs	
Available	89	(6%)	0	(0%)	6	(25%)	0	(0%)
In Use	1255	(85%)	53	(91%)	2	(8%)	6	(50%)
Down	132	(9%)	5	(9%)	16	(67%)	6	(50%)
Reserved	0	(0%)	0	(0%)	0	(0%)	0	(0%)
Total	1476		58		24		12	

Files, Directories and Paths

In Unix/Linux, everything is a file; even a directory is just a special file.

Each file on the system is identifiable by its absolute **path**.

A path is a string that specifies the name & location of a file (or directory); it consists of characters separated by zero or more **/** (“forward slash” character)

Examples of absolute paths:

/home/fit014

/usr/local/dos2unix/7.4.0/bin/dos2unix

/mnt/lustre/projects/fs19/tutorial/rawimages

are there **relative** paths?

Files, Directories and Paths (relative paths)

A relative path is one that **does not have** a leading / (“forward slash” character)

Each file on the system is identifiable by its absolute **path**. Any absolute path can also be specified as a **relative** path.

Examples of relative paths:

```
python          ?≡?  ./python
bin/dos2unix    ≡    ./bin/dos2unix
./fs19/tutorial/demo.zip
```

two special relative paths:

- . - this directory
- .. - the parent directory

Use of relative paths in your program **can break your access to data** if the program is run from another folder on the system.

Your personal and project spaces on MonARCH

`${HOME}`

- your personal `$HOME` directory
- good for source codes, configuration files, etc.
- not good for HPC with lots of data

`${HOME}/fs19`

- shortcut to the project folder for FIT3143
- a symbolic link to absolute path `/mnt/lustre/fs19`

`${HOME}/fs19/${USER}`

- your personal project space
- prepare and submit your jobs within this folder
- accessible only to you & your lecturer/tutors
- Note: `/mnt/lustre` sits on the fast parallel FS

Basic Navigation and Directory Management

- `ls` - list the file names in current directory
- `mkdir path` - create an empty directory
- `rmdir path` - removes an empty directory
- `cd path` - change into the given path/directory
 - `cd ..` - change into the enclosing directory
 - `cd` - change into your `$HOME` directory
- `tree` - visualise the directory structure

Or try this sequence:

```
cd; mkdir A; cd A; mkdir a b c; cd; tree
```

Something to try:

```
cd
mkdir 2019
ls
cd 2019
mkdir jan feb mar
ls
cd jan
mkdir wk1
cd ../mar
mkdir wk1
cd ../feb
mkdir wk1 wk2
cd
tree
```

Navigating your way within MonARCH

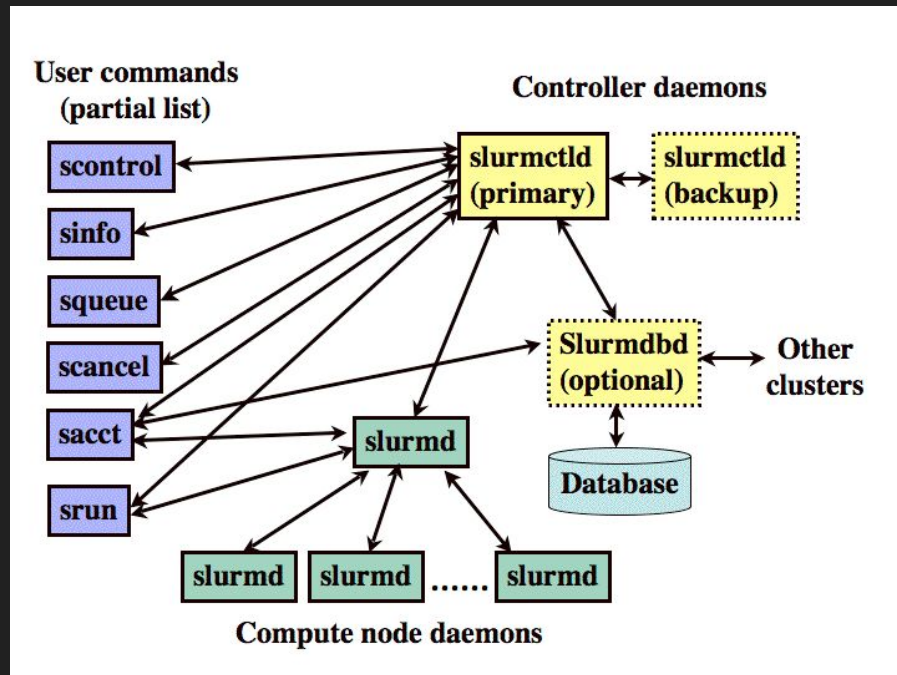
Key commands:

- `cat path` - displays the contents of *path* (try this on a folder?)
be careful running this on a binary file
- `cp src dest` - makes a copy of *src* into *dest* (both are paths)
- `mv src dest` - moves or renames files (or folders)
- `rm path` - permanently deletes the given file
- `ls -l` - long directory listing (equivalent shortcut `ll`)
- `cd` - go back `$HOME`
- `cd path` - change into the given path / folder

SLURM Scheduler

MonARCH uses SLURM as its resource manager and job scheduler:

- fair share scheduling policy
- no fixed CPU allocation, jobs are prioritised based on history of usage



Source: <https://slurm.schedmd.com/overview.html>

Interacting with the SLURM Job Scheduler

`sbatch jobscript.sh`

- to submit a job for batch execution;
- returns the unique JobID if successful

`squeue -u ${USER}`

- to display the status of my jobs on the queue

`scontrol show job {jobid}`

- to check the status of a job and other details

`sacct -j {jobid}`

- to check the accounting information of a given job

`scancel {jobid}`

- to delete a job that you own given its unique JobID

`scancel -u ${USER}`

- to delete all your jobs on the cluster

Commands unique to MonARCH:

`sq`

- this is a convenient shortcut to “`squeue -u ${USER}`”

`SQ`

- detailed listing of all user jobs (useful to determine job position in the queue)

`show_cluster`

- to see the status of all nodes on the cluster

Submitting your first sample job

1. Change to your project folder:

```
cd ~/fs19/${USER}
pwd
```

2. Make a copy of the demo.zip:

```
cp ../tutorial/demo.zip .
ll
```

3. Unzip the file and cd into demo:

```
unzip demo.zip
ll
cd demo
```

4. Submit the “demo.job”:

```
sbatch demo.job
sq
```

5. Check for output:

```
ll
tree
cat slurm*.out
```

6. View any one of the output JPEGs:

```
cd blackwhite
display filename
cd ../sharpened
```

MPI job script template (when running all 8 mpi processes on one node)

```
#!/bin/bash
#SBATCH --job-name=np_8_job      ### name your job
#SBATCH --time=00:30:00         ### hh:mm:ss or dd-hh:mm:ss
#SBATCH --mem=32G               ### set this to 8 ntasks x 4 GB
#SBATCH --ntasks=8              ### launch one python process
#SBATCH --cpus-per-task=1        ### single-threaded processes
#SBATCH --ntasks-per-node=8      ### set this to ntasks if ntasks <= 16
#SBATCH --account=fit3143        ### SLURM setting
#SBATCH --reservation=fit3143    ### reserved nodes
```

1/2 hours max per job

1 cpu max per task

Active as of 27/08/2020

```
module load openmpi/1.10.7-mlx
srun path-to-mpi-executable
```

one running max per user

MPI job script template *(same 8 processes but on two nodes)*

```
#!/bin/bash
```

```
#SBATCH --job-name=np_32_job
```

```
#SBATCH --time=00:30:00
```

```
#SBATCH --mem=32G
```

```
#SBATCH --ntasks=8
```

```
#SBATCH --cpus-per-task=1
```

```
#SBATCH --ntasks-per-node=4
```

```
#SBATCH --account=fit3143
```

```
#SBATCH --reservation=fit3143
```

```
### name your job
```

```
### hh:mm:ss or dd-hh:mm:ss
```

```
### memory setting is max @ 4 GB per core
```

```
### launch one python process
```

```
### single-threaded processes
```

```
### four per node
```

```
### SLURM setting
```

```
### reserved nodes
```

32 tasks is the current max per job

1/2 hours max per job

8 tasks for this job

1 cpu max per task

4 tasks per node, two nodes

Active as of 27/08/2020

```
module load openmpi/1.10.7-mlx  
srun path-to-mpi-executable
```

one running max per user

MPI job script template *(beyond 16 cores, you need two nodes)*

```
#!/bin/bash
```

```
#SBATCH --job-name=np_32_job
```

```
#SBATCH --time=00:30:00
```

```
#SBATCH --mem=32G
```

```
#SBATCH --ntasks=32
```

```
#SBATCH --cpus-per-task=1
```

```
#SBATCH --ntasks-per-node=16
```

```
#SBATCH --account=fit3143
```

```
#SBATCH --reservation=fit3143
```

```
### name your job
```

```
### hh:mm:ss or dd-hh:mm:ss
```

```
### memory setting is max @ 4 GB per core
```

```
### launch one python process
```

```
### single-threaded processes
```

```
### this is the max as each node has 16 cores only
```

```
### SLURM setting
```

```
### reserved nodes
```

32 tasks is the current max per job

1/2 hours max per job

32 tasks max per job

1 cpu max per task

Active as of 27/08/2020

```
module load openmpi/1.10.7-mlx
```

```
srun path-to-mpi-executable
```

one running max per user

MPI + OpenMP job script template (4 MPI processes on two nodes)

16 cpu cores allocated to this job, 8 cores per node

```
#!/bin/bash
```

```
#SBATCH --job-name=mpi_omp_4x4_job      ### name your job
```

```
#SBATCH --time=00:30:00                ### hh:mm:ss or dd-hh:mm:ss
```

1/2 hours max per job

```
#SBATCH --mem=32G                      ### memory setting is max @ 4 GB per core
```

```
#SBATCH --ntasks=4                    ### launch one python process
```

4 tasks for this job

```
#SBATCH --cpus-per-task=4              ### multi-threaded processes
```

4 cpu per mt process

```
#SBATCH --ntasks-per-node=2           ### two processes per node
```

2 tasks per node, two

```
#SBATCH --account=fit3143              ### SLURM setting
```

nodes

```
#SBATCH --reservation=fit3143          ### reserved nodes
```

Active as of 27/08/2020

```
module load openmpi/1.10.7-mlx
```

```
srun path-to-mpi-executable
```

one running max per user

MPI Demo Hello World Job

Self study and further reading

Install WinSCP or FileZilla and download/upload files into MonARCH

Running a text editor (**nedit** or **nano**). Alternatively, you may edit a file from Windows, upload to MonARCH, but will need to run: **dos2unix** to clean up the file.

Please read through the Getting-Started document:

https://docs.google.com/document/d/1Ewa49D00zYk1GK3Lg_ZxR-w0tUepW69U536w-owDZs

How to pick a secure but easy-to-remember password:

<https://xkcd.com/936/>

Further help with Unix/Linux:

<http://www.ee.surrey.ac.uk/Teaching/Unix/>

Thank you for your time.

Your questions welcome.

First point of contact \Rightarrow your tutor/lecturer

For help with cluster-related issues, email:

mcc-help@monash.edu

On the subject, please include this tag: **[FIT3143]**