

**FIT3143**

**LECTURE WEEK 8**

**PART I**

**INTRODUCTION TO INSTRUCTION  
LEVEL PARALLELISM (ILP)  
PROCESSORS**

## Overview

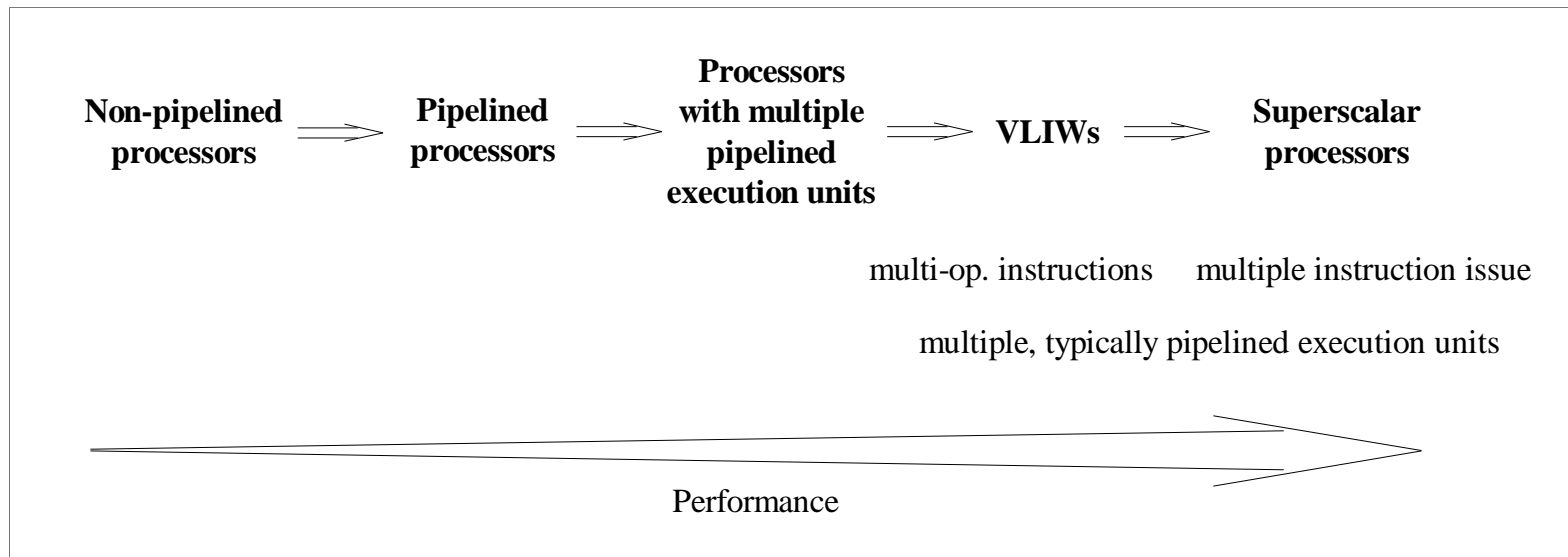
- Evolution and overview of ILP Processors
- Understanding of inter-instruction dependencies
  - Data, Control and Resource Dependencies
- What is instruction Scheduling?
- Sequential Consistency
- How fast can we go?

## Learning outcome(s) related to this topic

- Explain the fundamental principles of parallel computing architectures and algorithms (LO1)

## Evolution of ILP Processors

- Increasing processor performance has been provided by increased  
*Instruction Level Parallelism (ILP)*




## Instruction Level Parallelism

- Consider following code fragment

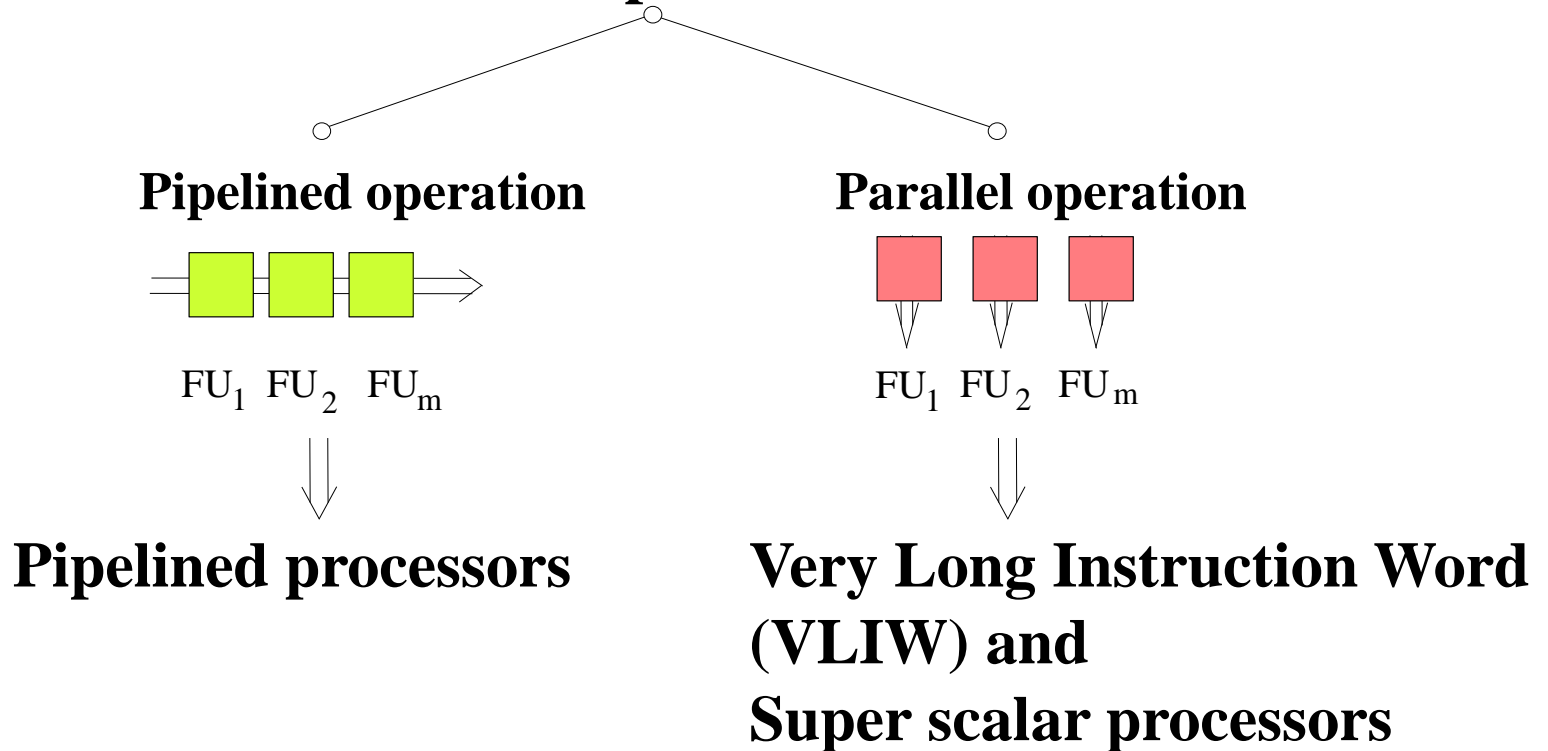
```
mul      r3, r4 → r8  
add      r1, r2 → r7  
load     r6, (r1)  
add      #5, r5
```

- Which instructions are unrelated?
  - Can be executed in any order

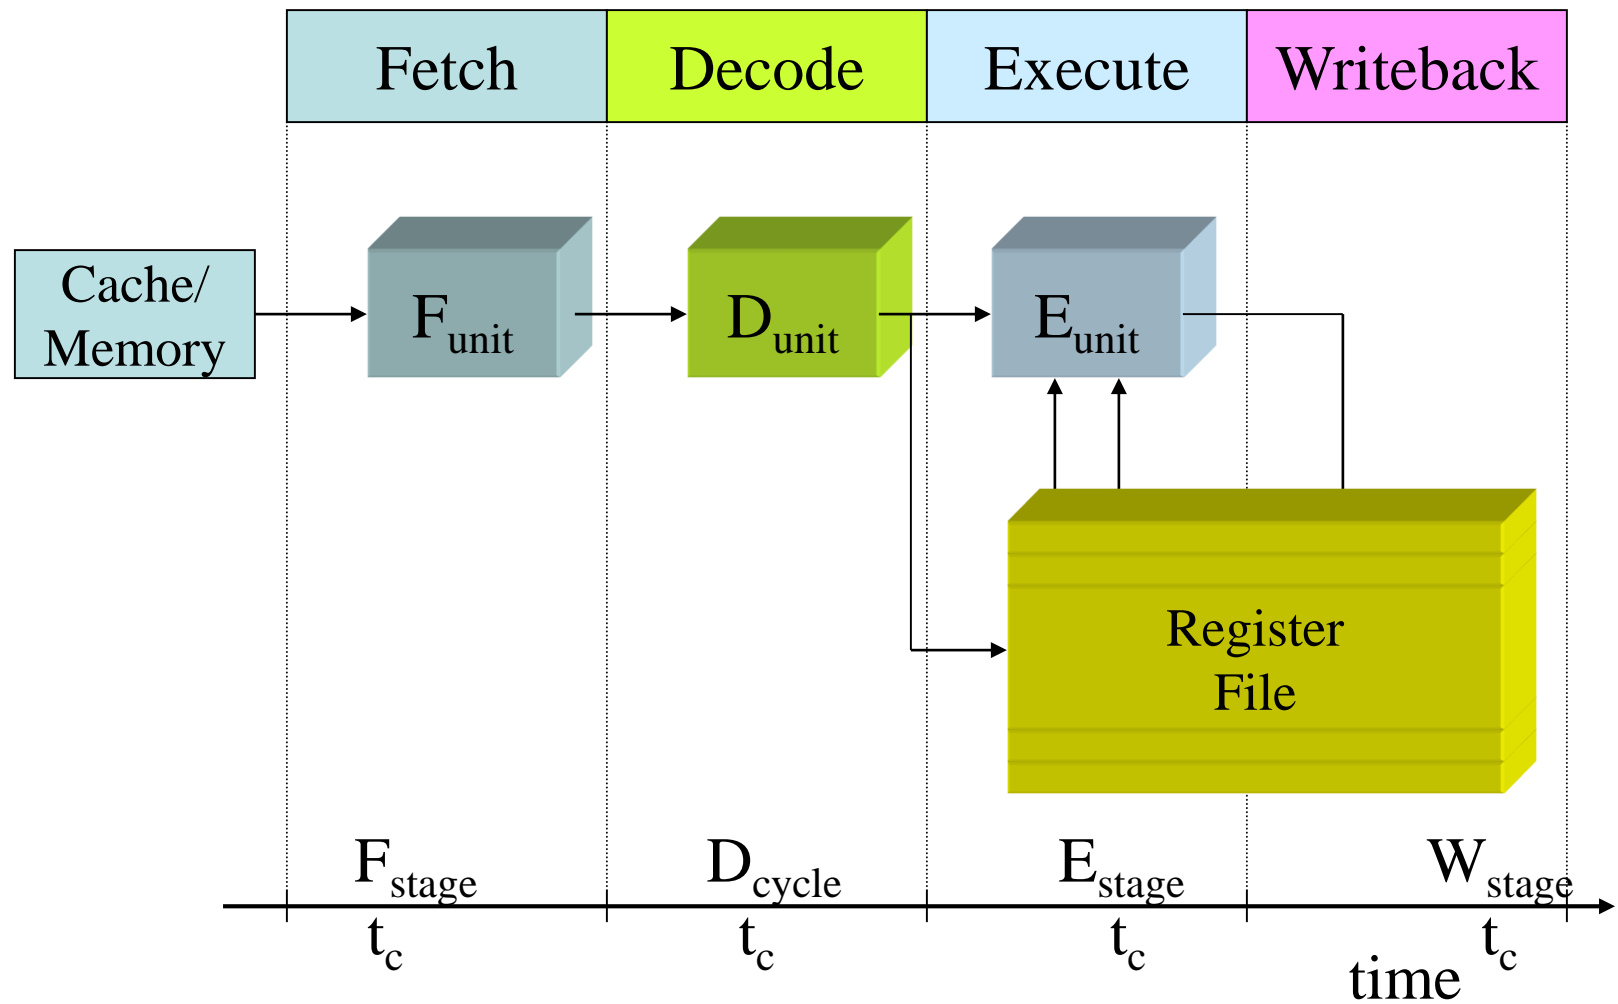


## Pipelining or parallel?

### Internal operating principle of IFP processors

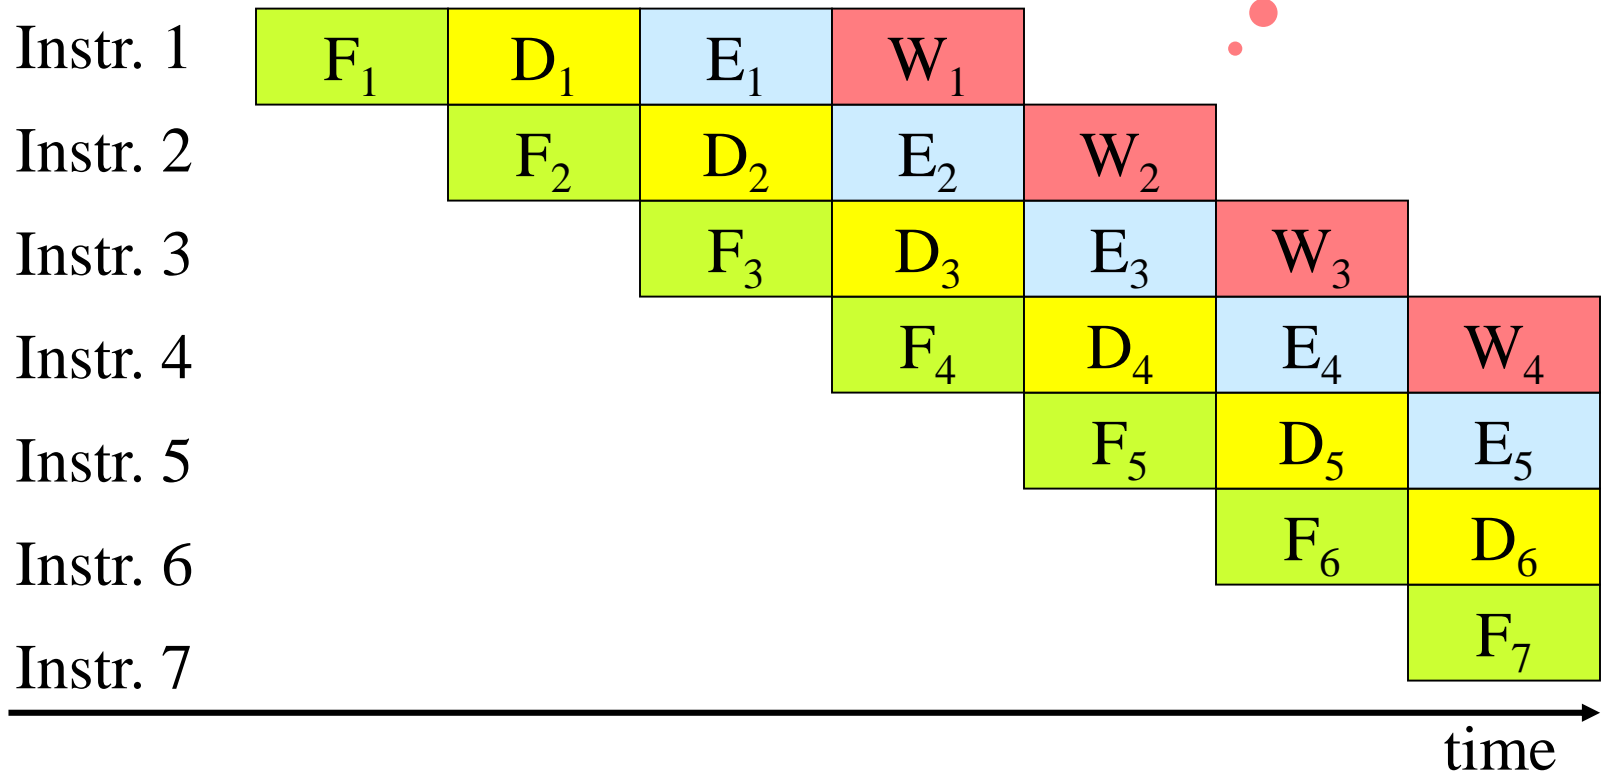


## Pipelining

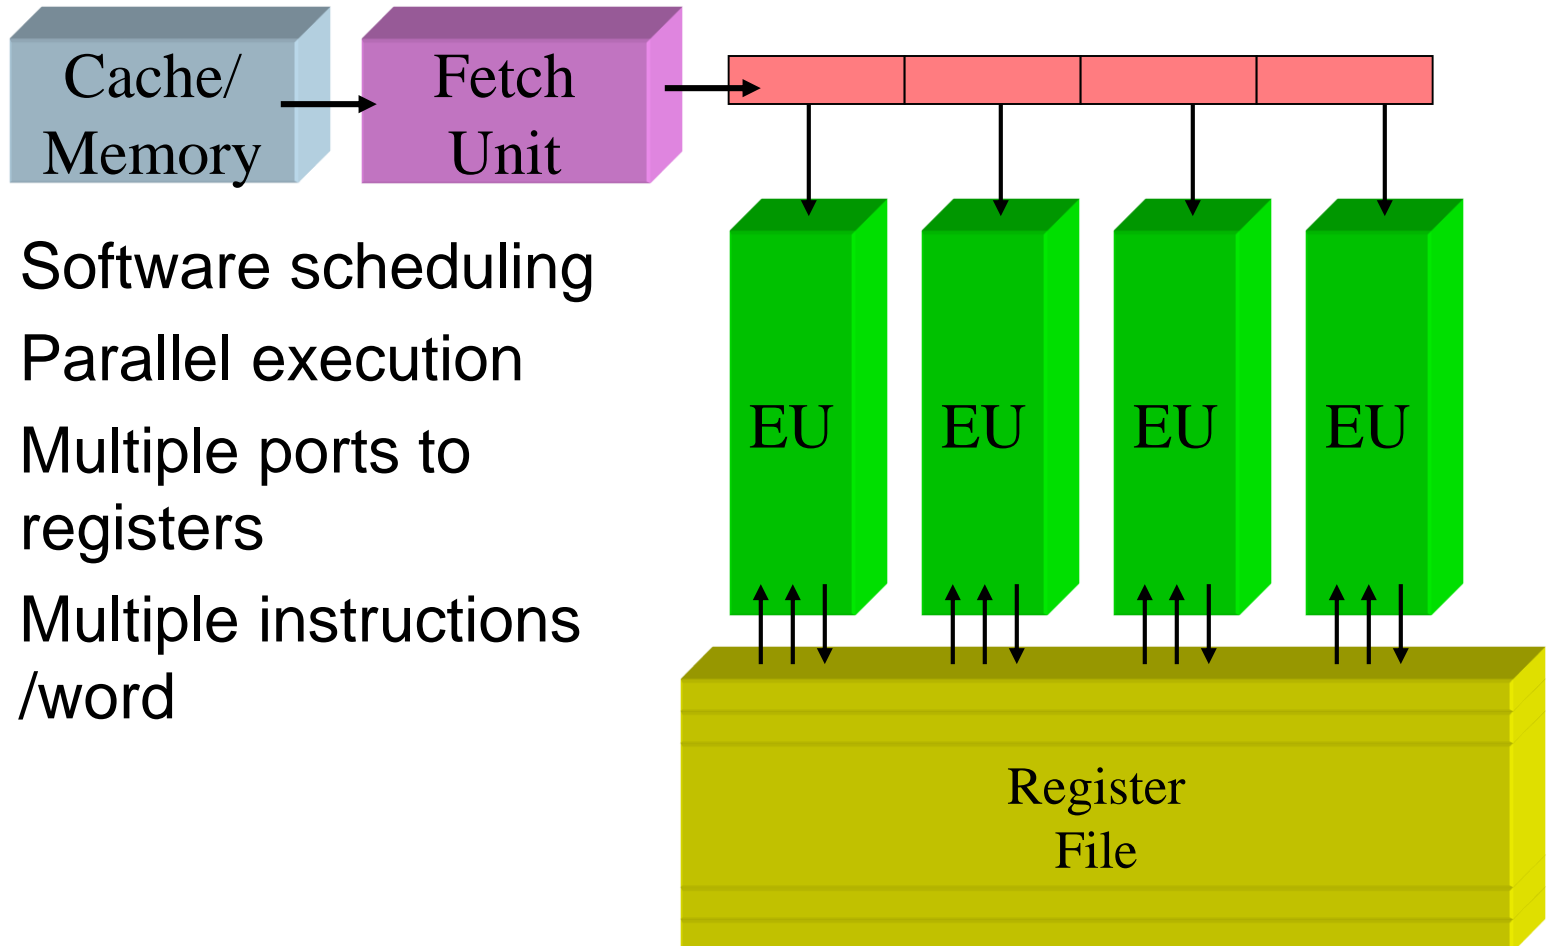


## Machine Utilisation with pipelining

Machine  
BUSY



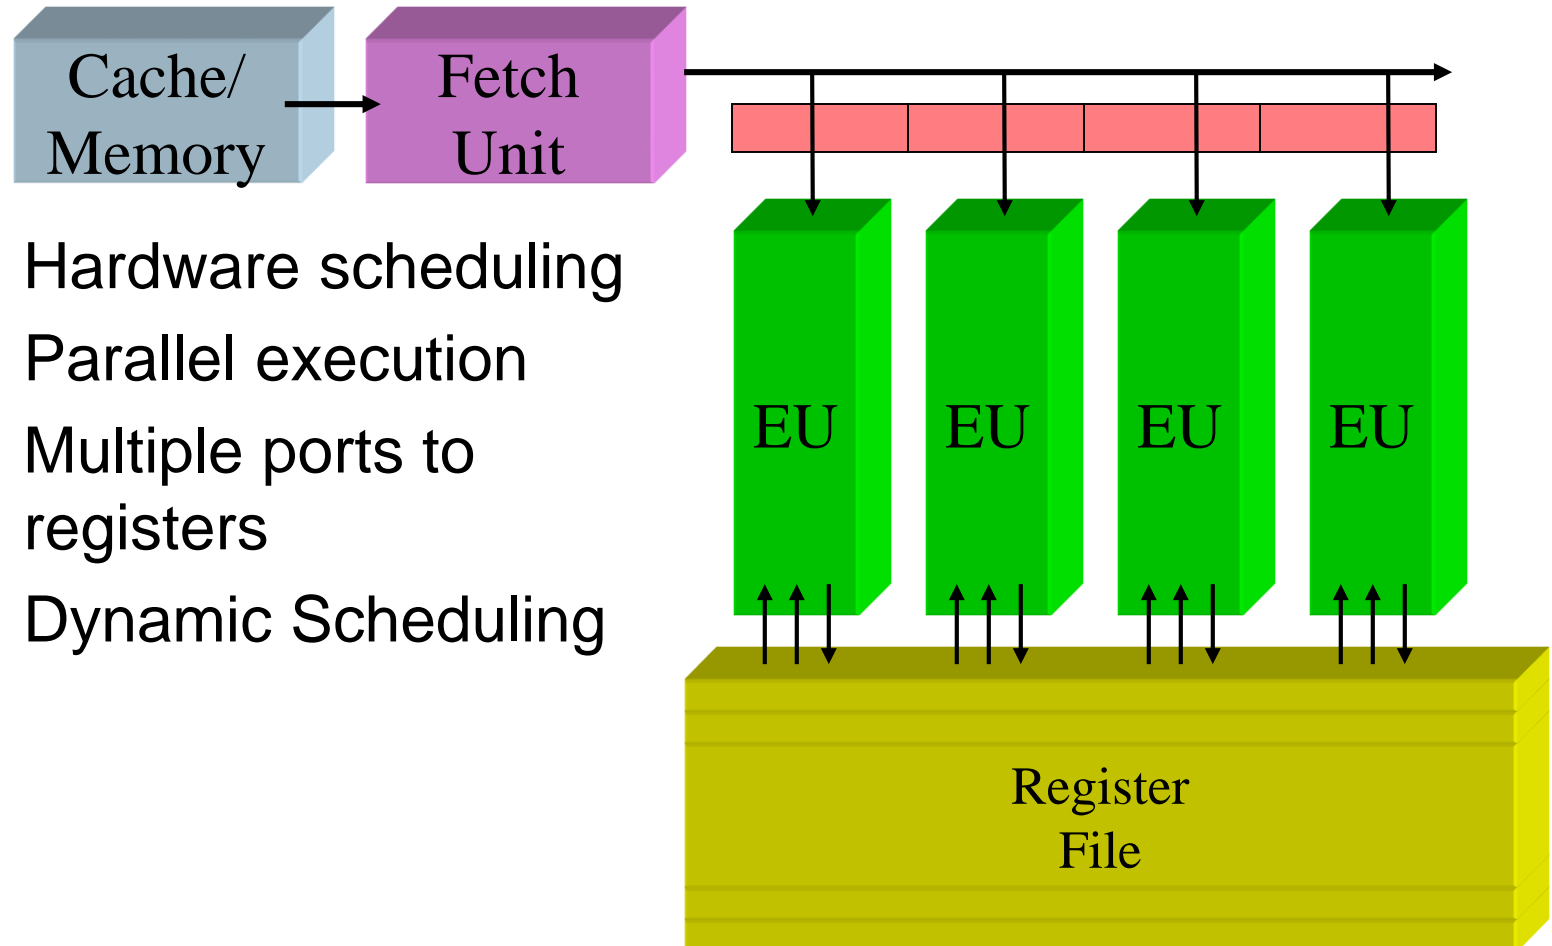
## VLIW Approach



- Software scheduling
- Parallel execution
- Multiple ports to registers
- Multiple instructions /word



## Superscalar Approach



- Hardware scheduling
- Parallel execution
- Multiple ports to registers
- Dynamic Scheduling

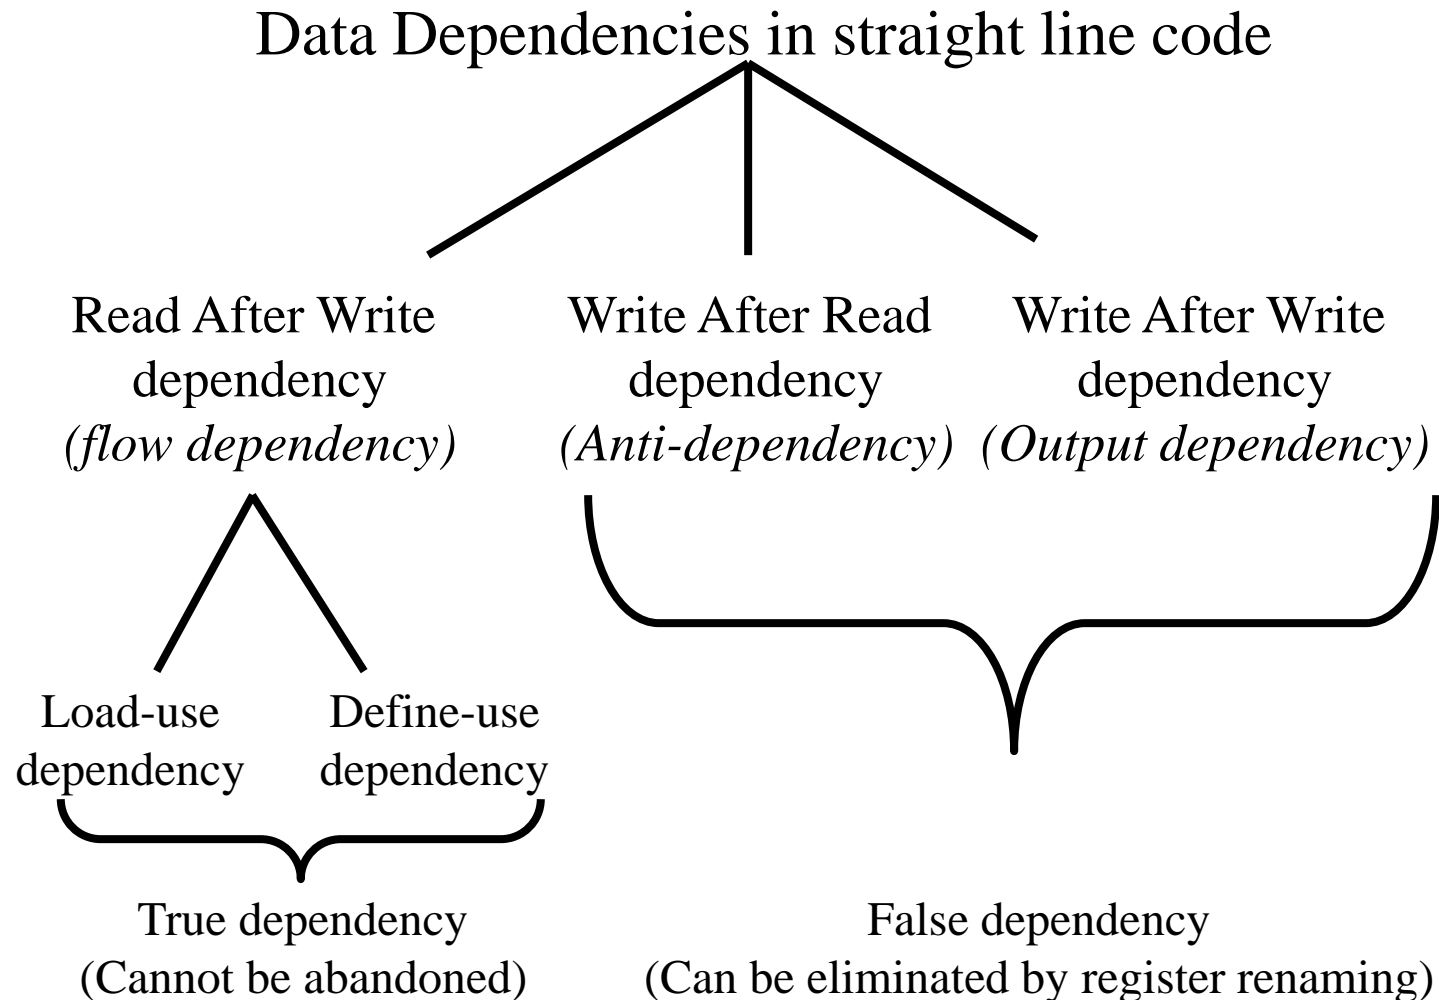
## Out-of-order Execution (OOOE)

- Look ahead in a window of instructions and find instructions that are *ready to execute*
  - Don't depend on data from previous instructions still not executed
  - Resources are available
- Out-of-order execution
  - Start instruction execution before execution of a previous instructions
- Advantages:
  - Help exploit Instruction Level Parallelism (ILP)
  - Help cover latencies (e.g., L1 data cache miss, divide)
- Can Compilers do the work ?
  - Compilers can *statically reschedule instructions*
  - Compilers do not have *run time information*
    - Conditional branch direction → limited to basic blocks
    - Data values, which may affect calculation time and control
    - Cache miss / hit

## **Dependencies between instructions**

- Data Dependencies
  - Future instructions depend on results of prior ones
  - Registers and Memory
- Control Dependencies
  - Branch dependencies
- Resource Dependencies
  - Number of ALUs, memory ports, etc

## Data Dependencies




## Data Flow Analysis


- Example

- (1) **r1**  $\leftarrow$  **r4** / **r7** ; assume divide takes 20 cycles
- (2) **r8**  $\leftarrow$  **r1** + **r2** ; **r1** depends on (1)
- (3) **r5**  $\leftarrow$  **r5** + 1 ; **r5** is independent
- (4) **r6**  $\leftarrow$  **r6** - **r3** ; **r6** & **r3** are independent
- (5) **r4**  $\leftarrow$  **r5** + **r6** ; **r5** depends on (3); **r6** depends on (4)
- (6) **r7**  $\leftarrow$  **r8** \* **r4** ; **r8** depends on (2); **r4** depends on (5)


Data Flow Graph



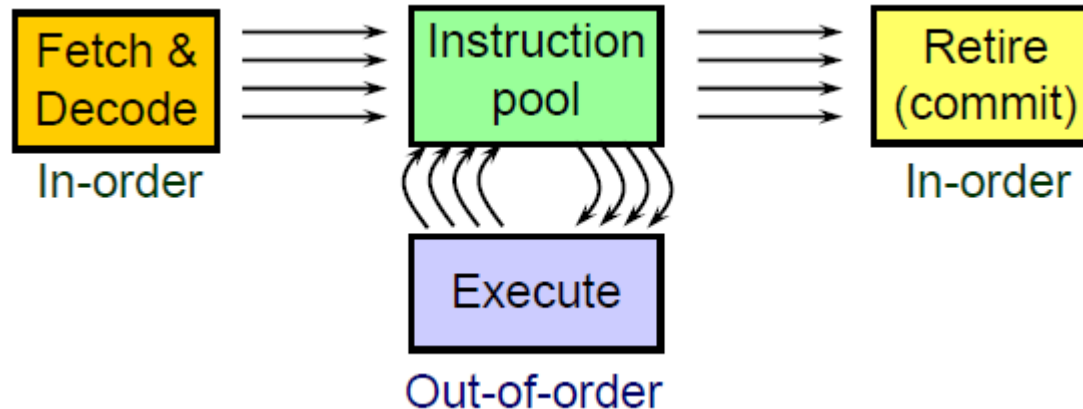
**In-order execution**



**Out-of-order execution**



## OOOE – General Scheme




- Fetch & decode instructions in parallel but in order, to fill instruction pool
- Execute ready instructions from the instructions pool
  - All the data required for the instruction is ready
  - Execution resources are available
- Once an instruction is executed
  - signal all dependant instructions that data is ready
- Commit instructions in parallel but in-order
  - Can commit an instruction only after all preceding instructions (in program order) have committed

## Example

- Assume that executing a divide operation takes 20 cycles

```
(1)  r1 ← r5 / r4
(2)  r3 ← r1 + r8
(3)  r8 ← r5 + 1
(4)  r3 ← r7 - 2
(5)  r6 ← r6 + r7
```



- Inst2 has a **RAW** dependency on r1 with Inst1
  - It cannot be executed in parallel with Inst1
- Can successive instructions pass Inst2 ?
  - Inst3 cannot since Inst2 must read r8 before Inst3 writes to it
  - Inst4 cannot since it must write to r3 after Inst2
  - Inst5 can

## False Dependencies

- OOOE creates new dependencies
  - **WAR**: write to a register which is read by an earlier inst.

(1)  $r3 \leftarrow r2 + r1$   
(2)  $r2 \leftarrow r4 + 3$

- **WAW**: write to a register which is written by an earlier inst.

(1)  $r3 \leftarrow r1 + r2$   
(2)  $r3 \leftarrow r4 + 3$


- These are *false dependencies*
  - There is no missing data
  - Still prevent executing instructions out-of-order
- Solution: *Register Renaming*



## Register Renaming

- Hold a pool of *physical registers*
- Map architectural registers into physical registers
  - Before an instruction can be sent for execution
    - **Allocate** a free physical register from a pool
    - The physical register **points** to the architectural register
  - When an instruction **writes** a result
    - Write the result value to the **physical register**
  - When an instruction needs data from a register
    - **Read** data from the **physical register** allocated to the latest inst which writes to the same arch register, and precedes the current instruction
    - If no such instruction exists, read directly from the arch register
  - When an instruction **commits**
    - **Move** the value from the **physical register** to the **arch register** it points


## Example



## Register Renaming Benefits

- Removes false dependencies
- Removes architecture limit for # of registers

## Mapping to Physical Registers



- $r1, r2, \dots, r6$  is referring to the architectural register while  $r1', r1''$  etc. are referring to physical registers. There can be multiple of them, depending on the processor.
- The results from each instruction will only be committed to the architectural registers in the in-order retirement stage.

## Executing Beyond Branches

- So far we do not look for instructions ready to execute beyond a branch
- Limited to the parallelism within a basic-block

- A basic-block is ~5 instruction long

(1) **r1**  $\leftarrow$  r4 / r7

(2) **r2**  $\leftarrow$  r2 + **r1**

(3) **r3**  $\leftarrow$  **r2** - 5

(4) **beq** **r3**, 0, 300      If the beq is predicted NT,

(5) **r8**  $\leftarrow$  r8 + 1      Inst 5 can be spec executed

- We would like to look beyond branches
  - But what if we execute an instruction beyond a branch and then it turns out that we predicted the wrong path ?
- Solution: *Speculative Execution*

## Speculative Execution

- Execution of instructions from a predicted (yet unsure) path
  - Eventually, path may turn wrong
- Implementation:
  - Hold a **pool** of all *not yet executed instructions*
    - Fetch instructions into the pool from a predicted path
  - Instructions for which all operands are **ready** *can be executed*
  - An instruction may change the processor state (commit) only when it is safe
    - An instruction commits only when all previous (in-order) instructions have committed **instructions commit** in-order
    - Instructions which follow a **branch commit** only after the branch commits
      - If a predicted branch is wrong all the instructions which follow it are flushed
- Register Renaming helps speculative execution
  - Renamed registers are **kept** until speculation is verified to be correct

## Example

		<u>cycle 1</u>	<u>cycle 2</u>
	(1) <b>r1</b> ← mem1	r1' ← mem1	
	(2) <b>WAW</b> r2 ← r2 + r1		r2' ← r2 + r1'
	(3) <b>r1</b> ← mem2	r1'' ← mem2	
	(4) <b>WAW</b> r3 ← r3 + r1		r3' ← r3 + r1''
	(5) jmp cond L2	predicted taken to L2	
Speculative Execution	(6) L2: <b>r1</b> ← mem3	r1''' ← mem3	
	(7) <b>WAR</b> r4 ← r5 + r1		r4' ← r5 + r1'''
	(8) <b>r5</b> ← 2	r5' ← 2	
	(9) r6 ← r5 + 2		r6' ← r5' + 2

- Instructions 6-9 are speculatively executed
  - If the prediction turns **wrong**, they will be **flushed**
- If the branch was **predicted taken**
  - The instructions from the other path would have been speculatively executed

## Data Dependencies in Loops

- Not all data dependencies are obvious by examining the source code.
- Consider
  - do  $l = 2, n$ 
    - $X(l) = A * X(l-1) + B$
- Recurrence between iteration  $l$  and  $l-1$
- Compiler needs to analyze loop expressions to detect and handle recurrences correctly.

## Control Dependencies

- Consider the following code

```
mul    r1, r2 → r3
```

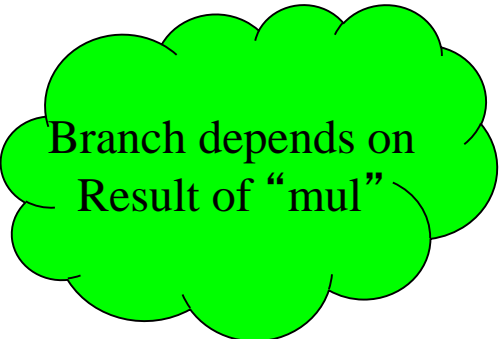
```
jz     zproc
```

```
sub r4, r1 → r1
```

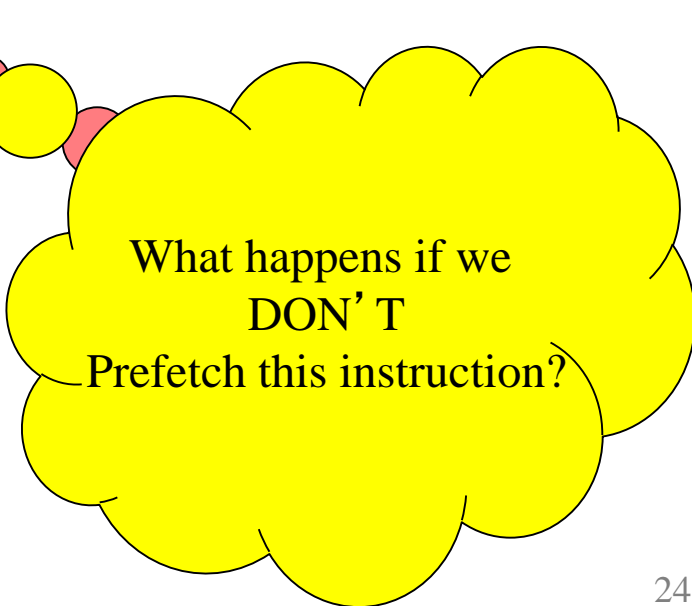
```
:
```

```
zproc: load r1, x
```

```
:
```



Branch depends on  
Result of “mul”



What happens if we  
DON' T  
Prefetch this instruction?



## Is this a problem?

- Ratio of branches
  - General Purpose 22 % to 39%
  - Scientific 5% - 11 %
- Conditional to unconditional
  - General Purpose 46% to 83%
  - Scientific 53 % to 83 %
- For 25% 1 in 4 instruction is a branch
- Effect on pipeline?



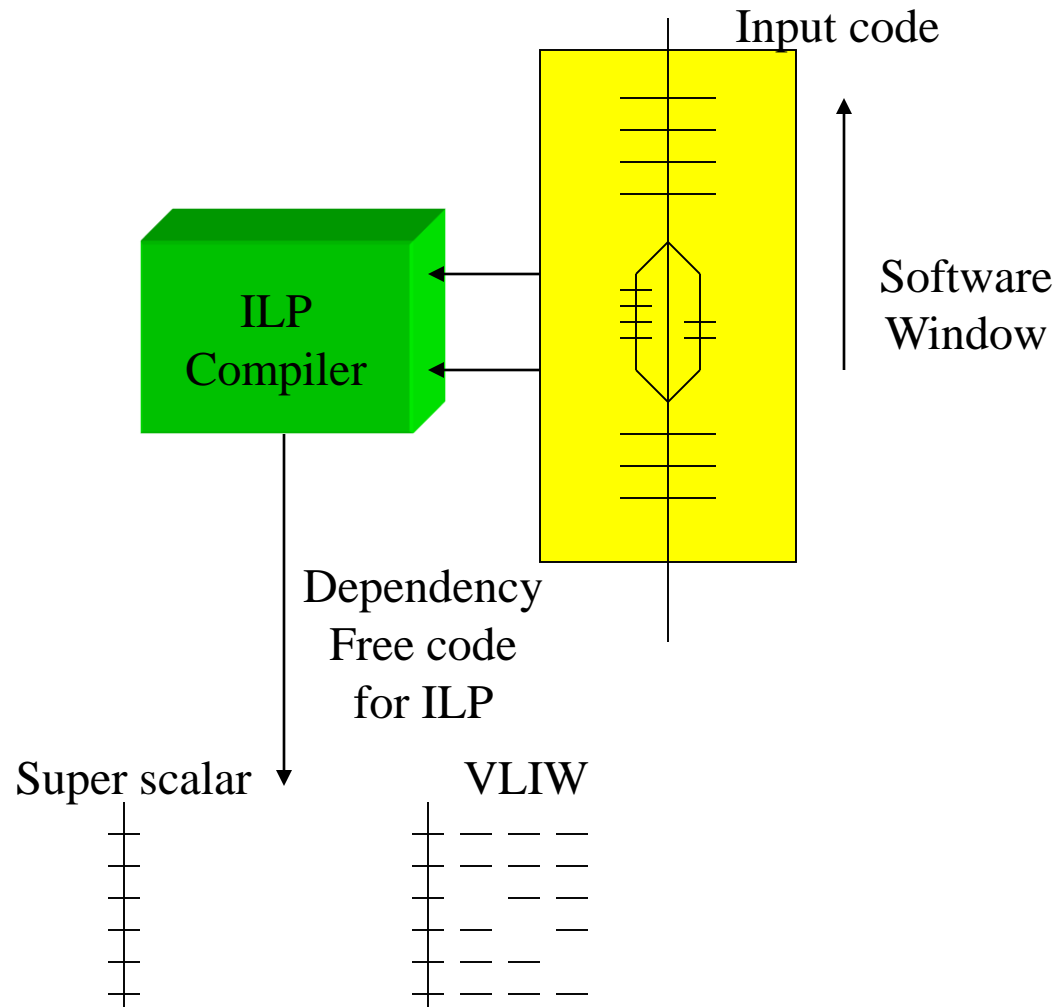
## Resource Dependencies

- Consider the following code
  - div            r1, r2  $\rightarrow$  r3
  - div            r4, r2  $\rightarrow$  r5
- Insufficient dividers to issue both instructions in parallel
- Possible resources
  - buses, execution units, buffers (storage slots)

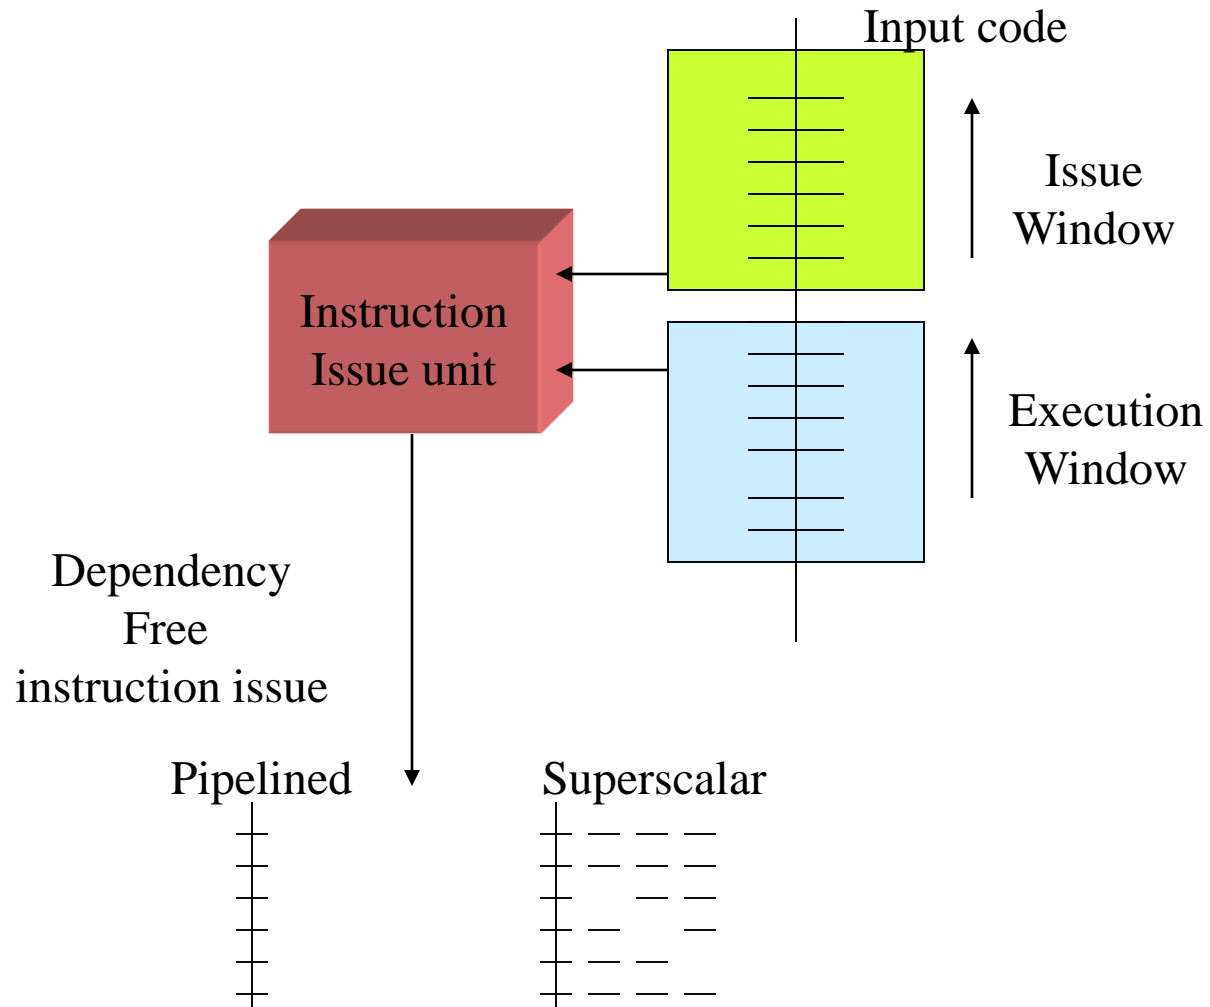
## Instruction scheduling

- When instructions are processed (that is, issued, renamed, executed and so forth) in parallel, it is often necessary to detect and resolve dependencies between instructions.
- Three main approaches
  - Static scheduling
  - Dynamic scheduling
  - Hybrid

## Static vs Dynamic scheduling



## Static vs Dynamic scheduling ...



## ILP Instruction Scheduling

- Static Scheduling boosted by parallel code optimisation
  - Performed entirely by the compiler
  - Processor receives dependency-free and optimized code for parallel execution
  - *Typical for VLIW and a few pipelined processors (e.g. MIPS)*

## ILP Instruction Scheduling

- Dynamic Scheduling without static parallel code optimisation
  - Performed entirely by the processor
  - The code is not optimised for parallel execution. The Processor detects and resolves dependencies on its own
  - *Early ILP processors (e.g. CDC 6600, IBM 360/91)*



## ILP Instruction Scheduling

- Dynamic scheduling boosted by static parallel code optimisation
  - Performed by the processor in conjunction with the parallel optimizing compiler
  - The processor receives optimised code for parallel execution, but it detects and resolves dependencies of its own
  - *Usual practice for pipelined superscalar processors (e.g. RS/6000, PL.8/XL, i960 and QTC SF960)*

## Preserving Sequential Consistency

- When instructions are executed in parallel, processor must be careful to preserve the sequential consistency.  

div r1, r2 → r3

add r5, r6 → r7

jz anywhere
- Must make sure the *jz* instruction uses the condition codes set by the *add*, not the *div* which might take longer.

## Types of consistency

- Strong consistency
  - Preserves the actual execution order
- Weak consistency
  - Produces correct result
  - Can execute out of order providing the code still delivers the correct result

## How fast can we go?

- Performance is limited by
  - underlying algorithm (dependencies)
  - compiled code (false dependencies, staging)
  - actual hardware (resource restrictions)
- Study all three to work out maximum speedup

## How fast can we go?

### Earliest methods to achieve speed ups for scientific programs

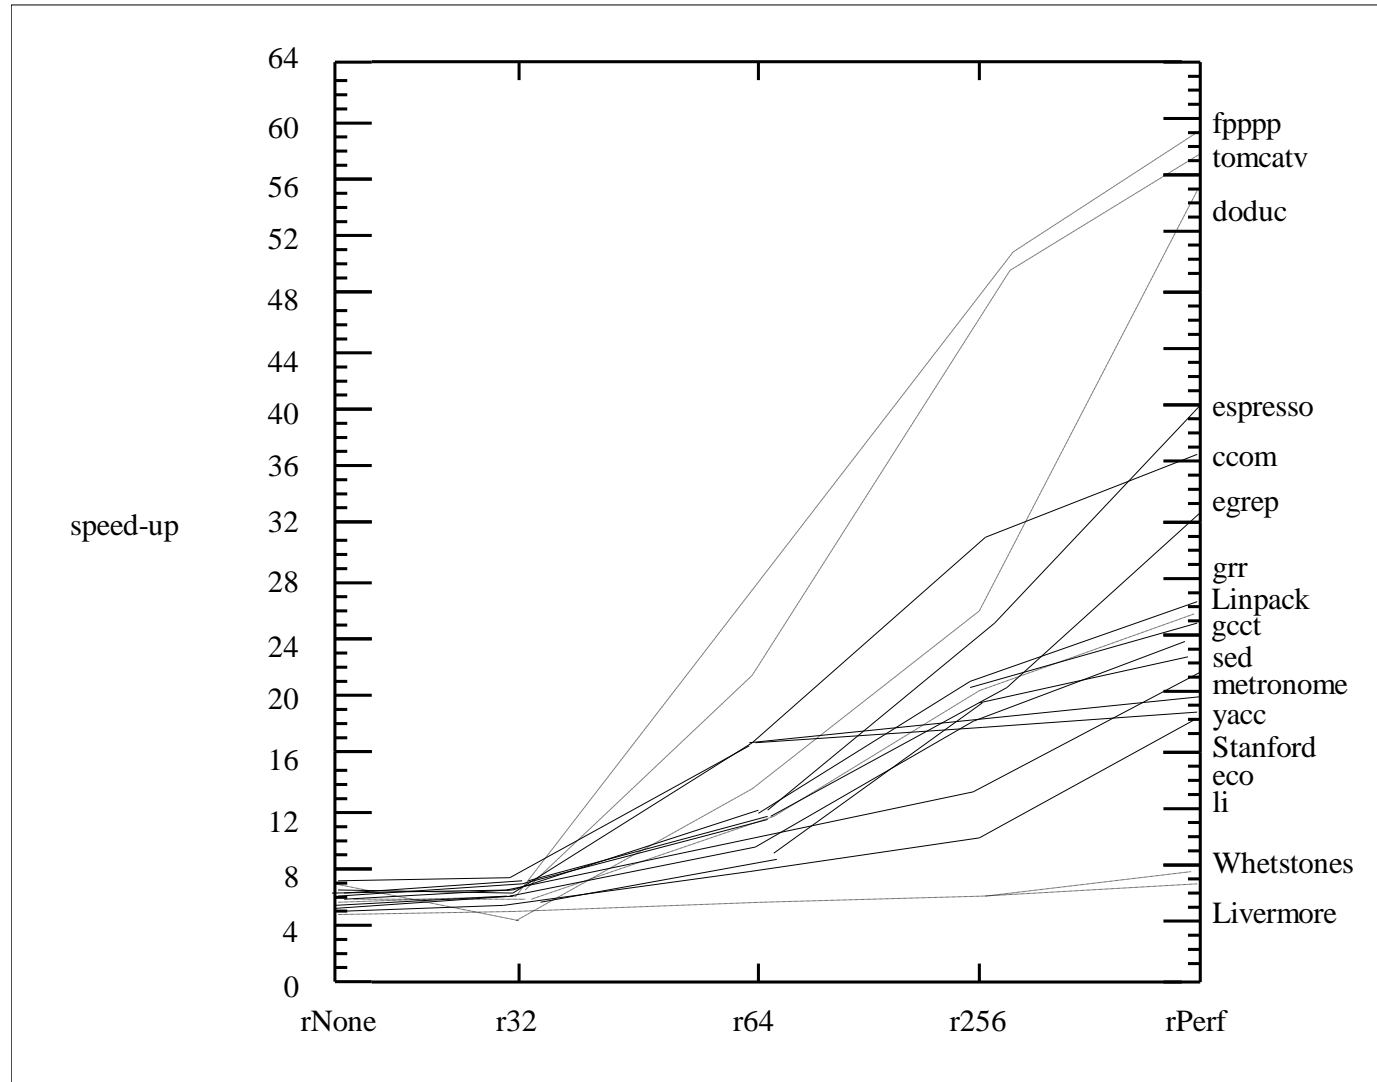
Paper	Benchmark	Speed-up for general programs		Speed-up for scientific programs	
		range	average	range	average
Tjaden, Flynn 1970	31 Library programs			1.2 - 3.2	<b>1.9</b>
Kuck & at. 1972	20 FORTRAN programs			1.2 - 17	<b>### 4</b>
Riseman, Foster 1972	7 FORTRAN/ assembly programs	1.2 - 3.0	<b>1.8</b>	1.4 - 1.6	<b>1.6</b>
Jouppy 1989	8 Modula-2 programs	1.6 - 2.2	<b>1.9</b>	2.4 - 3.3	<b>2.8</b>
Lam, Wilson 1992	6 SPECmarks + 4 others	1.5 - 2.8	<b>2.1</b>	2 - 293	

## How fast can we go?

### Earliest methods to achieve speed ups for scientific programs

Paper	Benchmark	Speed-up for general purpose programs		Speed-up for scientific programs	
		range	average	range	average
Riseman, Foster 1972	7 FORTRAN/assembly programs	8 - 100	<b>42</b>	30 - 120	<b>75</b>
Nicolaus, Fisher 1984	22 numerical programs			3 - 988	<b>12</b>
Kumar 1988	4 FORTRAN programs			475 - 3500	<b>1839</b>
Butler & at.1991	9 SPEC benchmarks	38 - 200	<b>170</b>	17 - 1165	<b>509</b>
Wall 1991	6SPEC benchmarks, 13 others	16 -41	<b>28</b>	57 - 60	<b>59</b>
		7 - 37	<b>24</b>	6 - 27	<b>18</b>
Austin, Sohi 1992	SPEC benchmarks	13 - 942	<b>288</b>	51 - 33749	<b>6188</b>
Lam, Wilson 1992	6SPEC benchmarks, 4 others	174 - 3283	<b>1400</b>	844 - 188 470	<b>64500</b>
		47 - 265	<b>229</b>		

## Effect of register renaming



## Summary

- Evolution and overview of ILP Processors
  - Pipelines, VLIW or super scaler
- Understanding of inter-instruction dependencies
  - Data, Control and Resource Dependencies
- Instruction Scheduling
  - Static & Dynamic scheduling
- Sequential Consistency
  - When instructions are executed in parallel, processor must be careful to preserve the sequential consistency.
- How fast can we go?
  - Performance is limited by underlying algorithm compiled code and actual hardware (resource restrictions)



## References

Sima .D, Fountain .T, Karsuk .P, “Chapter 4: Introduction to ILP-Processors”, in *Advanced Computer Architectures – A Design Space Approach*, Addison Wesley Longman, 1997, pp. 113 – 135.